

# Graph Neural Networks Designs

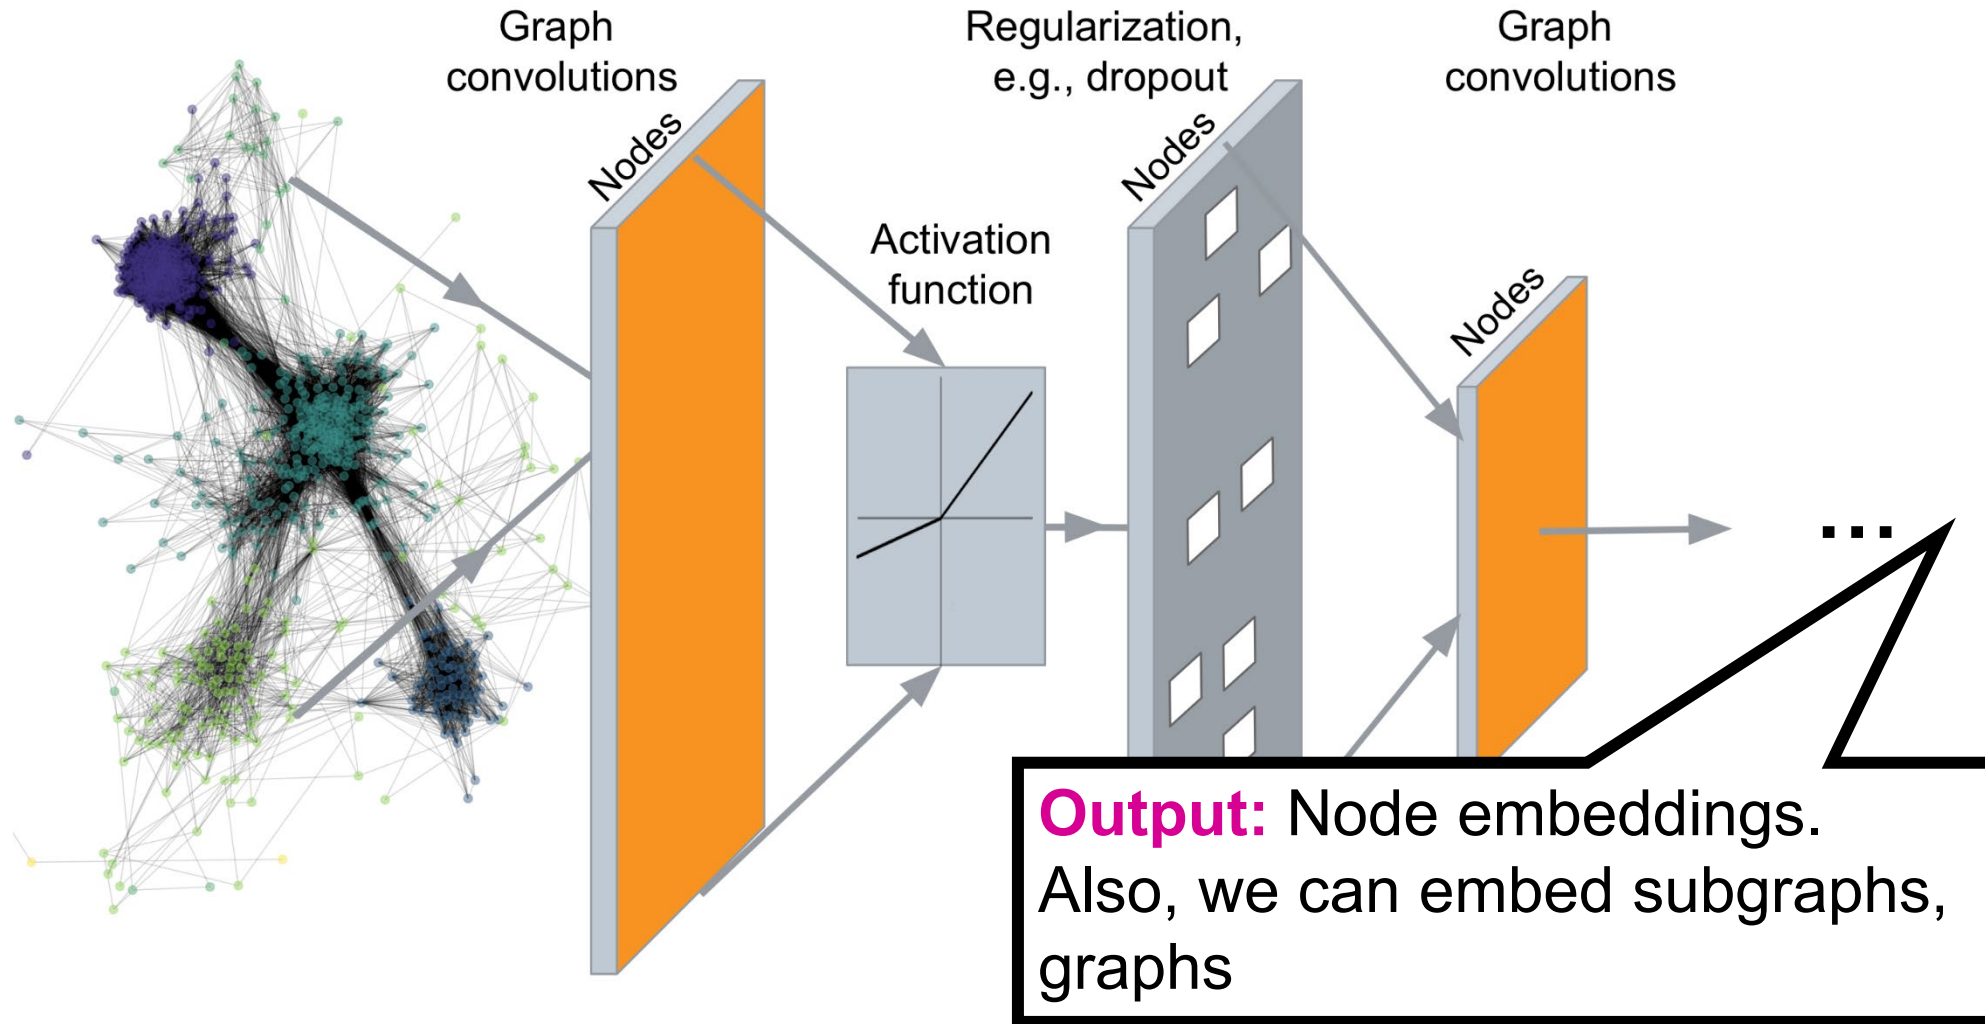
CPSC483: Deep Learning on Graph-Structured Data

Rex Ying

# Readings


- Readings are updated on the website (syllabus page)
- **Lecture 3 readings:**
  - [Graph representation learning: methods and application](#)
  - [Inductive Representation Learning for Large Graphs \(GraphSAGE\)](#)
- **Lecture 4 readings:**
  - [Semi-Supervised Classification with Graph Convolutional Networks](#)
  - [Principled Neighborhood Aggregation on Graph Nets](#)

# Recap: Deep Graph Encoders




# Recap: Graph Convolutional Networks

- **Idea:** Node's neighborhood defines a computation graph



Determine node  
computation graph




Propagate and  
transform information

**Learn how to propagate information across the graph to compute node features**

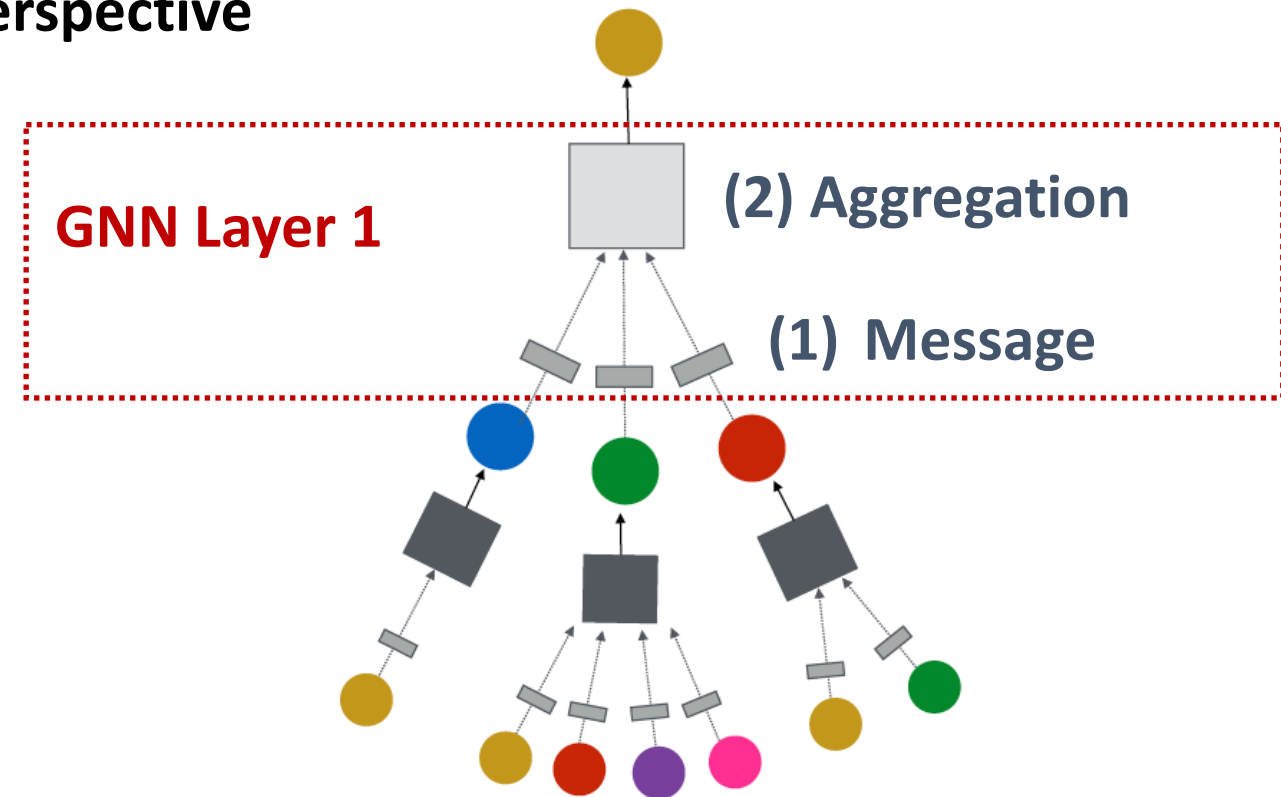
# Recap: Aggregate Neighbors (1)

- **Intuition:** Nodes aggregate information from their neighbors using neural networks



# A General GNN Framework (1)


- **GNN Layer = Message + Aggregation**
  - Different instantiations under this perspective
  - GCN, GraphSAGE, GAT, ...



# A General GNN Framework (2)


- **Connect GNN layers into a GNN**
- Stack layers sequentially
- Ways of adding skip connections

(3) Layer connectivity



# A General GNN Framework (3)

- **Idea: Raw input graph  $\neq$  computational graph**
  - Graph feature augmentation
  - Graph structure augmentation




(4) Graph augmentation



# A General GNN Framework (4)

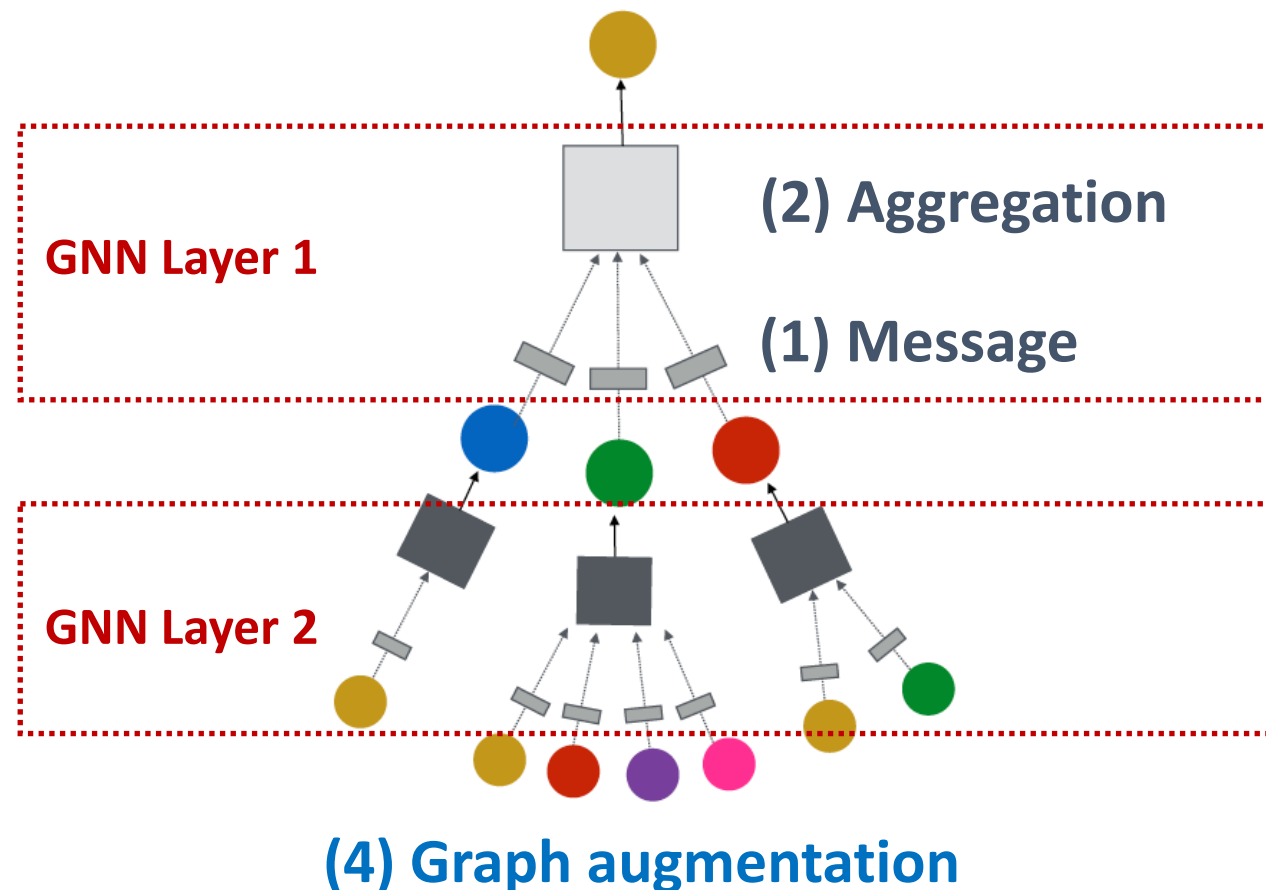
- **How do we train a GNN**
  - Supervised/Unsupervised objectives
  - Node/Edge/Graph level objectives



(4) Graph augmentation

# A General GNN Framework (5)

(5) Learning objective



# Content

- **A Single Layer of a GNN**
- **Stacking Layers of a GNN**
- **Graph Manipulation in GNNs**

# Outline of Today's Lecture



- **A Single Layer of a GNN**
- Stacking Layers of a GNN
- Graph Manipulation in GNNs

# A Single layer of a GNN

# A Single GNN Layer: Two Steps

- **Idea of a GNN Layer:**

- Compress a set of vectors into a single vector
- **Two step process:**
  - (1) Message
  - (2) Aggregation




# Message Computation

- **(1) Message computation**

## Message function

- **Intuition:** Each node will create a message, which will be sent to other nodes later
- **Example:** A Linear layer  $\mathbf{m}_u^{(l)} = \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}$ 
  - Multiply node features with weight matrix  $\mathbf{W}^{(l)}$



# Message Aggregation


- **(2) Aggregation**

- **Intuition:** Each node will aggregate the messages from node  $v$ 's neighbors

$$\mathbf{h}_v^{(l)} = \text{AGG}^{(l)} \left( \left\{ \mathbf{m}_u^{(l)}, u \in N(v) \right\} \right)$$

- **Example:** Sum( $\cdot$ ), Mean( $\cdot$ ) or Max( $\cdot$ ) aggregator

$$\mathbf{h}_v^{(l)} = \text{Sum}(\{\mathbf{m}_u^{(l)}, u \in N(v)\})$$





# Message Aggregation: Issue

- **Issue:** Information from node  $v$  itself **could get lost**

- Computation of  $\mathbf{h}_v^{(l)}$  does not directly depend on  $\mathbf{h}_v^{(l-1)}$

- **Solution:** Include  $\mathbf{h}_v^{(l-1)}$  when computing  $\mathbf{h}_v^{(l)}$

- **(1) Message:** compute message from node  $v$  itself

- Usually, a **different message computation** will be performed

$$\text{●} \text{●} \text{●} \quad \mathbf{m}_u^{(l)} = \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)} \qquad \text{●} \quad \mathbf{m}_v^{(l)} = \mathbf{B}^{(l)} \mathbf{h}_v^{(l-1)}$$

- **(2) Aggregation:** After aggregating from neighbors, we can **aggregate the message from node  $v$  itself**

- Via **concatenation** or **summation** **Then aggregate from node itself**

$$\mathbf{h}_v^{(l)} = \text{CONCAT} \left( \text{AGG} \left( \left\{ \mathbf{m}_u^{(l)}, u \in N(v) \right\} \right) \right) \mathbf{m}_v^{(l)}$$

First aggregate from neighbors

# A Single GNN Layer: Message and Aggregation

- **Putting things together:**

- **(1) Message:** each node computes a message


$$\mathbf{m}_u^{(l)} = \text{MSG}^{(l)} \left( \mathbf{h}_u^{(l-1)} \right), u \in \{N(v) \cup v\}$$

- **(2) Aggregation:** aggregate messages from neighbors

$$\mathbf{h}_v^{(l)} = \text{AGG}^{(l)} \left( \left\{ \mathbf{m}_u^{(l)}, u \in N(v) \right\}, \mathbf{m}_v^{(l)} \right)$$

- **Nonlinearity (activation):** Adds expressiveness

- Often written as  $\sigma(\cdot)$ :  $\text{ReLU}(\cdot)$ ,  $\text{Sigmoid}(\cdot)$ , ...
- Can be added to **message or aggregation**

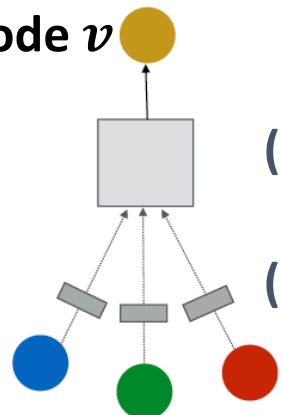


# Classical GNN Layers: GCN (1)

- **(1) Graph Convolutional Networks (GCN)**

$$\mathbf{h}_v^{(l)} = \sigma \left( \mathbf{W}^{(l)} \sum_{u \in N(v)} \frac{\mathbf{h}_u^{(l-1)}}{\sqrt{|N(u)| |N(v)|}} \right)$$

- **How to write this as Message + Aggregation?**

$$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N(v)} \underbrace{\mathbf{W}^{(l)}}_{\text{Aggregation}} \underbrace{\frac{\mathbf{h}_u^{(l-1)}}{\sqrt{|N(u)| |N(v)|}}}_{\text{Message}} \right)$$


Node  $v$


(2) Aggregation

(1) Message

# Classical GNN Layers: GCN (2)

- **Graph Convolutional Networks (GCN)**

$$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N(v)} \mathbf{w}^{(l)} \frac{\mathbf{h}_u^{(l-1)}}{\sqrt{|N(u)||N(v)|}} \right)$$



- **Message:**

- Each Neighbor:  $\mathbf{m}_u^{(l)} = \frac{1}{\sqrt{|N(u)||N(v)|}} \mathbf{w}^{(l)} \mathbf{h}_u^{(l-1)}$

- **Aggregation:**

- **Sum** over messages from neighbors, then apply activation

- $\mathbf{h}_v^{(l)} = \sigma \left( \text{Sum} \left( \left\{ \mathbf{m}_u^{(l)}, u \in N(v) \right\} \right) \right)$

# Classical GNN Layers: GraphSAGE

- **GraphSAGE**

$$\mathbf{h}_v^{(l)} = \sigma \left( \mathbf{W}^{(l)} \cdot \text{CONCAT} \left( \mathbf{h}_v^{(l-1)}, \text{AGG} \left( \left\{ \mathbf{h}_u^{(l-1)}, \forall u \in N(v) \right\} \right) \right) \right)$$

- **How to write this as Message + Aggregation?**

- **Message** is computed within the  $\text{AGG}(\cdot)$

- **Two-stage aggregation**

- **Stage 1:** Aggregate from node neighbors

$$\mathbf{h}_{N(v)}^{(l)} \leftarrow \text{AGG} \left( \left\{ \mathbf{h}_u^{(l-1)}, \forall u \in N(v) \right\} \right)$$

- **Stage 2:** Further aggregate over the node itself

$$\mathbf{h}_v^{(l)} \leftarrow \sigma \left( \mathbf{W}^{(l)} \cdot \text{CONCAT}(\mathbf{h}_v^{(l-1)}, \mathbf{h}_{N(v)}^{(l)}) \right)$$

# GraphSAGE Neighborhood Aggregation

- **Mean:** Take a weighted average of neighbors

$$\text{AGG} = \underbrace{\sum_{u \in N(v)} \mathbf{h}_u^{(l-1)}}_{\text{Aggregation}} \underbrace{|N(v)|}_{\text{Message computation}}$$

- **Pool:** Transform neighbor vectors and apply symmetric vector function  $\text{Mean}(\cdot)$  or  $\text{Max}(\cdot)$

$$\text{AGG} = \underbrace{\text{Mean}}_{\text{Aggregation}}(\underbrace{\{\text{MLP}(\mathbf{h}_u^{(l-1)}), \forall u \in N(v)\}}_{\text{Message computation}})$$

- **LSTM:** Apply LSTM to reshuffled of neighbors (not order invariant)

$$\text{AGG} = \underbrace{\text{LSTM}}_{\text{Aggregation}}([\mathbf{h}_u^{(l-1)}, \forall u \in \pi(N(v))])$$

# GraphSAGE: L2 Normalization


- **$\ell_2$  Normalization:**

- **Optional:** Apply  $\ell_2$  normalization to  $\mathbf{h}_v^{(l)}$  at every layer
- $\mathbf{h}_v^{(l)} \leftarrow \frac{\mathbf{h}_v^{(l)}}{\|\mathbf{h}_v^{(l)}\|_2} \quad \forall v \in V$  where  $\|u\|_2 = \sqrt{\sum_i u_i^2}$  ( $\ell_2$ -norm)
- Without  $\ell_2$  normalization, the embedding vectors have different scales ( $\ell_2$ -norm) for vectors
- In some cases (not always), normalization of embedding results in performance improvement
- **Benefit:** more efficient retrieval (finding the nearest embedding vector) through locality sensitive hashing (**LSH**)

# GNN Layer in Practice (1)

- In practice, these classic GNN layers are a great starting point
  - We can often get better performance by considering a general GNN layer design
  - Concretely, we can include modern deep learning modules that proved to be useful in many domains

A suggested GNN Layer





# GNN Layer in Practice Overview

- **Overview: many modern deep learning modules can be incorporated into a GNN layer**

- **Batch Normalization**

- Normalize embeddings in a minibatch
- Stabilize neural network training

- **Layer Normalization (alternative to BN)**

- Normalize individual output embeddings


- **Dropout**

- Prevent overfitting

- **More:**

- Any other useful deep learning modules

## A suggested GNN Layer



# Batch Normalization

- **Goal:** Stabilize neural networks training
- **Idea:** Given a batch of inputs (node embeddings)
  - Re-center the node embeddings into zero mean
  - Re-scale the variance into unit variance

**Input:**  $\mathbf{X} \in \mathbb{R}^{N \times D}$   
 $N$  node embeddings

**Trainable Parameters:**

$\gamma, \beta \in \mathbb{R}^D$

**Output:**  $\mathbf{Y} \in \mathbb{R}^{N \times D}$   
Normalized node embeddings

**Step 1:**  
**Compute the mean and variance over  $N$  embeddings**


$$\mu_j = \frac{1}{N} \sum_{i=1}^N \mathbf{x}_{i,j}$$
$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}_{i,j} - \mu_j)^2$$

**Step 2:**  
**Normalize the feature using computed mean and variance**

$$\hat{\mathbf{x}}_{i,j} = \frac{\mathbf{x}_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$
$$\mathbf{y}_{i,j} = \gamma_j \hat{\mathbf{x}}_{i,j} + \beta_j$$


# Dropout

- **Goal:** Regularize a neural net to prevent overfitting.
  - Dropout ratio: percentage of dimensions that are dropped out
- **Idea:**
  - **During training:** with some probability  $p$ , randomly set neurons to zero (turn off)
  - **During testing:** Use all the neurons for computation
  - In PyTorch, use `model.train()` and `model.eval()` to toggle



# Dropout for GNNs

- In GNN, Dropout is applied to **the linear layer in the message function**
  - A simple message function with linear layer:  $\mathbf{m}_u^{(l)} = \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}$



# Activation (Non-linearity)

Apply activation to  $i$ -th dimension of embedding  $\mathbf{x}$

- **Rectified linear unit (ReLU)**

$$\text{ReLU}(\mathbf{x}_i) = \max(\mathbf{x}_i, 0)$$

- Commonly used

- **Sigmoid**

$$\sigma(\mathbf{x}_i) = \frac{1}{1 + e^{-\mathbf{x}_i}}$$




- Used only when you want to restrict the range of your embeddings

- **Parametric ReLU**

$$\text{PReLU}(\mathbf{x}_i) = \max(\mathbf{x}_i, 0) + a_i \min(\mathbf{x}_i, 0)$$


$a_i$  is a **trainable parameter**

- Sometimes performs better than ReLU
- See [Pytorch documentation](#) for more non-linearity functions



# GNN Layer in Practice

- **Summary:** Various deep learning modules can be included into a GNN layer for better performance
- **Designing novel GNN layers is still an active research frontier!**
- **Suggested resources:** You can explore diverse GNN designs or try out your own ideas in [GraphGym](#)



# Outline of Today's Lecture

- A Single Layer of a GNN
- **Stacking Layers of a GNN**
- Graph Manipulation in GNNs


# Stacking Layers of a GNN



# Stacking GNN Layers (1)


- **How to connect GNN layers into a GNN?**
- Stack layers sequentially
- Ways of adding skip connections

(3) Layer connectivity



# Stacking GNN Layers (2)

- **How to construct a Graph Neural Network?**
  - **The standard way:** Stack GNN layers sequentially
  - **Input:** Initial raw node feature  $\mathbf{x}_v$
  - **Output:** Node embeddings  $\mathbf{h}_v^{(L)}$  after  $L$  GNN layers

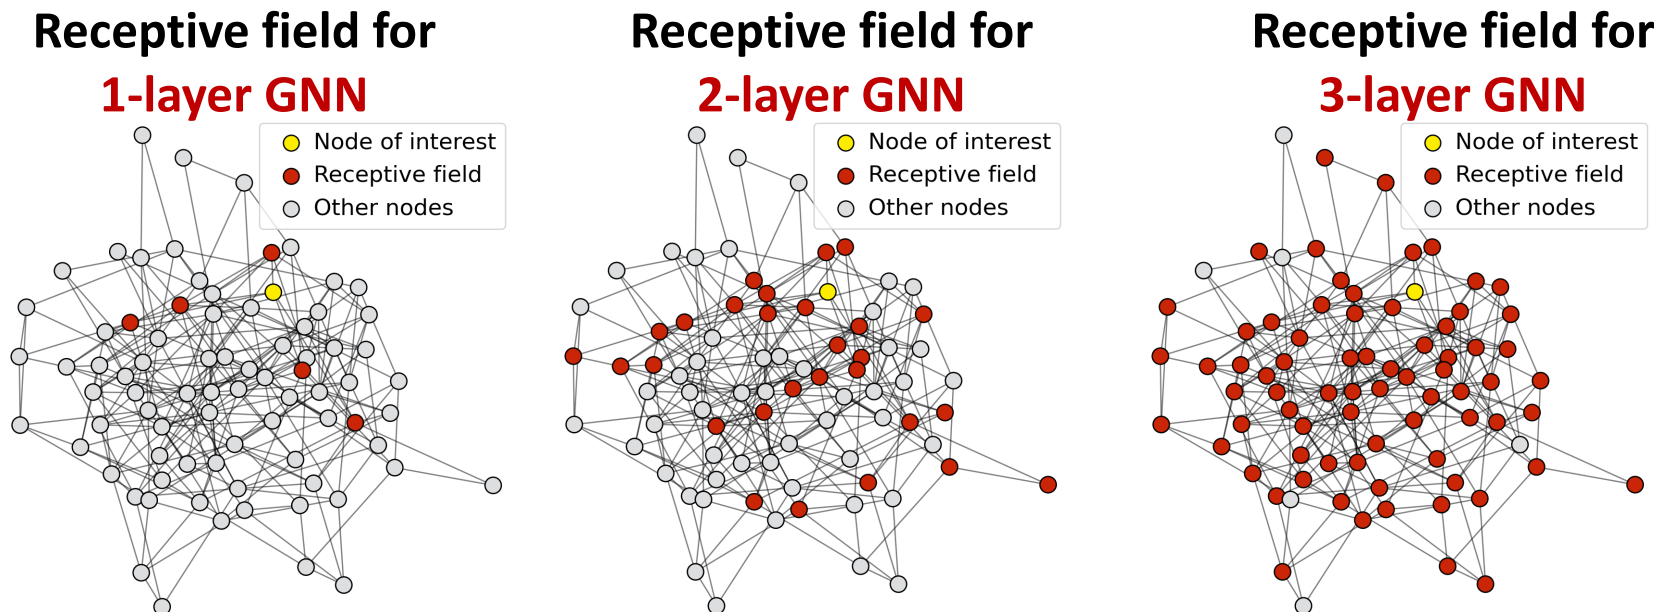


# The Over-smoothing Problem

- **The Issue of stacking many GNN layers**
  - GNN suffers from **the over-smoothing problem**
- **The over-smoothing problem:** all the node embeddings converge to the same value
  - This is bad because we **want to use node embeddings to differentiate nodes**
- **Why does the over-smoothing problem happen?**

# Receptive Field of a GNN (1)

- **Receptive field:** the set of nodes that determine the embedding of a node of interest
  - In a  $K$ -layer GNN, each node has a receptive field of  $K$ -hop neighborhood



# Receptive Field of a GNN (2)

- **We can explain over-smoothing via the notion of receptive field**
  - We knew **the embedding of a node is determined by its receptive field**
    - If two nodes have highly-overlapped receptive fields, then their embeddings are highly similar
  - **Stack many GNN layers** → **nodes will have highly-overlapped receptive fields** → **node embeddings will be highly similar** → **suffer from the over-smoothing problem**
- **Next:** how do we overcome over-smoothing problem?


# Design GNN Layer Connectivity

- **What do we learn from the over-smoothing problem?**
- **Lesson 1: Be cautious when adding GNN layers**
  - Unlike neural networks in other domains (CNN for image classification), **adding more GNN layers do not always help**
  - **Step 1: Analyze the necessary receptive field** to solve your problem. E.g., by computing the diameter of the graph
  - **Step 2:** Set number of GNN layers  $L$  to be a bit more than the receptive field we like. Tune it as a hyper-parameter.
- **Question:** How to enhance the expressive power of a GNN, **if the number of GNN layers is small?**

# Expressive Power for Shallow GNNs (1)


- How to make GNN more expressive without more message passing layers?
- **Solution 1:** Increase the expressive power **within each GNN layer**
  - In our previous examples, each transformation or aggregation function only include one linear layer
  - We can make aggregation / transformation become a deep neural network!

If needed, each box could include a **3-layer MLP**



# Expressive Power for Shallow GNNs (2)

- **How to make GNN more expressive without more message passing layers?**
- **Solution 2:** Add layers that do not pass messages
  - A GNN does not necessarily only contain GNN layers. e.g., we can add **MLP layers** (applied to each node) before and after GNN layers, as **pre-process layers** and **post-process layers**



**Pre-processing layers:** Important when encoding node features is necessary.

E.g., when nodes represent images/text

**Post-processing layers:** Important when reasoning / transformation over node embeddings are needed


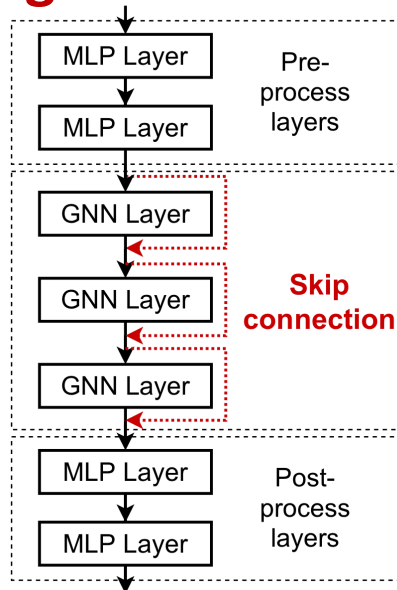
E.g., graph classification, knowledge graphs

**In practice, adding these layers works great!**



# Design GNN Layer Connectivity

- What if my problem still requires many GNN layers?
- **Lesson 2: Add skip connections in GNNs**
  - **Observation from over-smoothing:** Node embeddings in earlier GNN layers can sometimes better differentiate nodes
  - **Solution:** We can increase the impact of earlier layers on the final node embeddings, **by adding shortcuts in GNN**



## Idea of skip connections:

Before adding shortcuts:

$$F(x)$$

After adding shortcuts:

$$F(x) + x$$


(also called **residual stream** in the context of Transformer)

# Idea of Skip Connections

- **Why do skip connections work?**

- **Intuition:** Skip connections create **a mixture of models**
- $N$  skip connections  $\rightarrow 2^N$  possible paths
- Each path could have up to  $N$  modules
- We automatically get **a mixture of shallow GNNs and deep GNNs**

Path 2: skip this module




Path 1: include this module

(a) Conventional 3-block residual network

=

All the possible paths:  
 $2 * 2 * 2 = 2^3 = 8$



(b) Unraveled view of (a)

# Example: GCN with Skip Connections

- A standard GCN layer

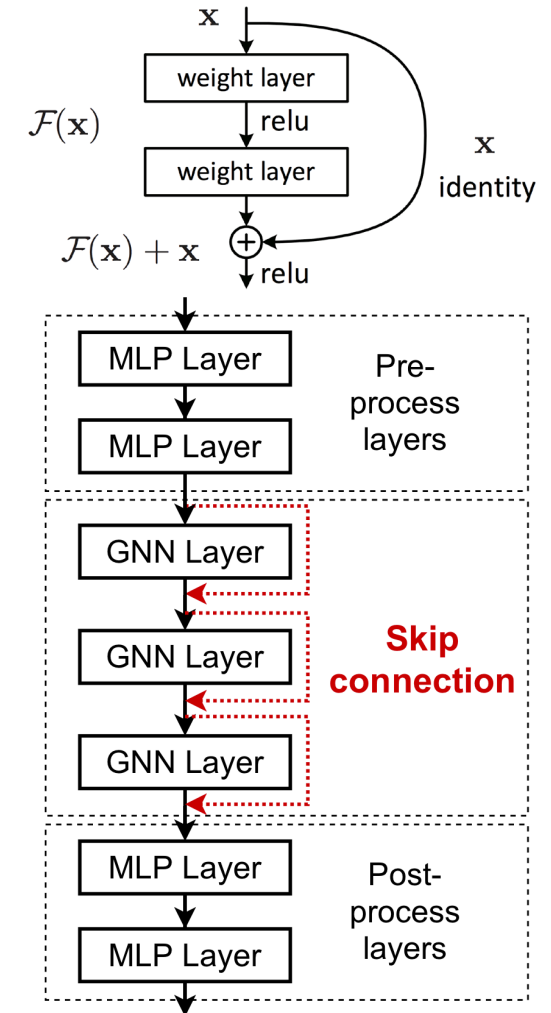
$$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N(v)} \mathbf{W}^{(l)} \frac{\mathbf{h}_u^{(l-1)}}{|N(v)|} \right)$$

This is our  $F(\mathbf{x})$

- A GCN layer with skip connection

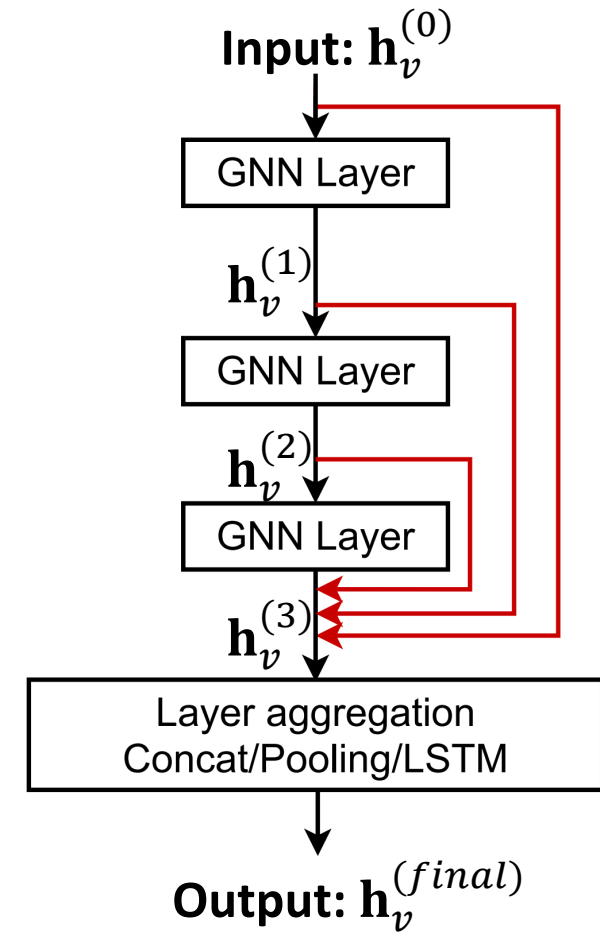
$$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N(v)} \mathbf{W}^{(l)} \frac{\mathbf{h}_u^{(l-1)}}{|N(v)|} + \mathbf{h}_v^{(l-1)} \right)$$

$F(\mathbf{x})$                       +                       $\mathbf{x}$



# Other Options of Skip Connections

- **Other options:** Directly skip to the last layer
  - The final layer directly **aggregates from the all the node embeddings** in the previous layers




# Outline of Today's Lecture

- A Single Layer of a GNN
- Stacking Layers of a GNN
- **Graph Manipulation in GNNs**

# Graph Manipulation in GNNs

# Graph GNN Framework

- **Idea: Raw input graph  $\neq$  computational graph**
  - Graph feature augmentation
  - Graph structure manipulation



(4) Graph manipulation

# Why Manipulating Graphs

Our assumption so far has been

- **Raw input graph = computational graph**
  - **Feature level**
    - The input graph **lacks features**
  - **Structure level:**
    - The graph is **too sparse** → insufficient message passing
    - The graph is **too dense** → message passing is too costly
    - The graph is **too large** → cannot fit the computational graph into a GPU
- It's **unlikely that the input graph happens to be the optimal computation graph** for embeddings



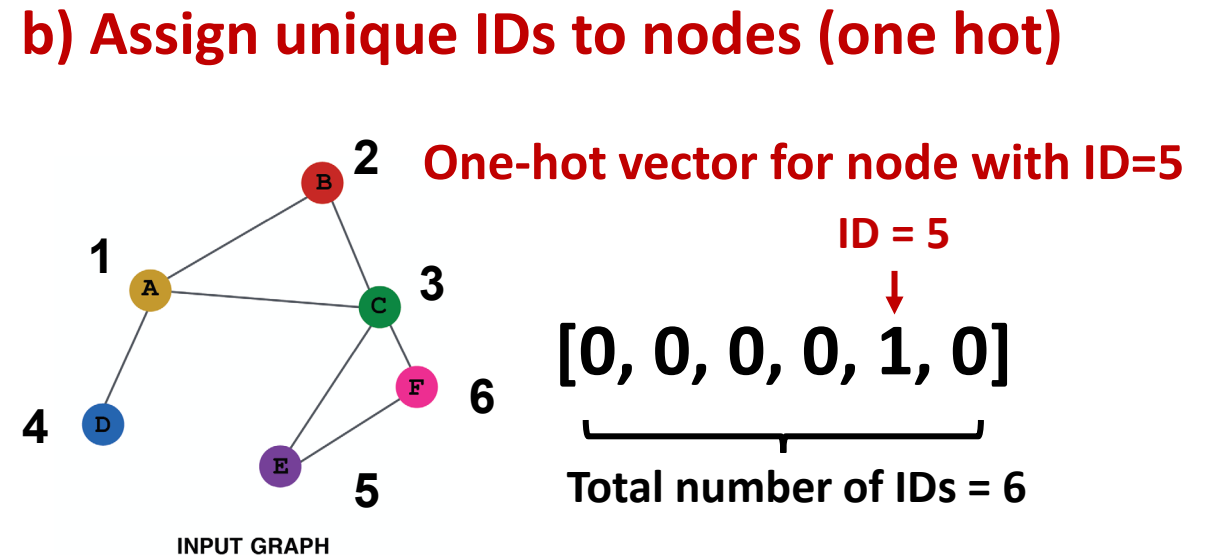

# Graph Manipulation Approaches

- **Graph Feature manipulation**
  - The input graph **lacks features** → **feature augmentation**
- **Graph Structure manipulation**
  - The graph is **too sparse** → **Add virtual nodes / edges**
  - The graph is **too dense** → **Sample neighbors when doing message passing**
  - The graph is **too large** → **Sample subgraphs to compute embeddings**
    - Will cover later in lecture: Scaling up GNNs

# Feature Augmentation on Graphs (1)

## Why do we need feature augmentation?

- **(1) Input graph does not have node features**
  - This is common when we only have the adj. matrix
- **Standard approaches:**
  - **a) Assign constant values to nodes**
  - **b) Assign unique IDs to nodes (one hot)**



# Feature Augmentation on Graphs (2)

- Feature augmentation: **constant** vs. **one-hot**


	Constant node feature	One-hot node feature
Expressive power	<b>Medium</b> . All the nodes are identical, but GNN can still learn from the graph structure	<b>High</b> . Each node has a unique ID, so node-specific information can be stored
Inductive learning (Generalize to unseen nodes)	<b>High</b> . Simple to generalize to new nodes: we assign constant feature to them, then apply our GNN	<b>Low</b> . Cannot generalize to new nodes: new nodes introduce new IDs, GNN doesn't know how to embed unseen IDs
Computational cost	<b>Low</b> . Only 1 dimensional feature	<b>High</b> . $O( V )$ dimensional feature, cannot apply to large graphs
Use cases	Any graph, inductive settings (generalize to new nodes)	Small graph, transductive settings (no new nodes)

# Feature Augmentation on Graphs (3)


## Why do we need feature augmentation?

- (2) Certain structures are hard to learn by GNN
- **Example:** Cycle count feature
  - Can GNN learn the length of a cycle that  $v_1$  resides in?
  - **Unfortunately, no** (we will talk more about it in Lecture 9)

$v_1$  resides in a cycle with length 3




$v_1$  resides in a cycle with length 4




# Feature Augmentation on Graphs (4)

- $v_1$  cannot differentiate which graph it resides in
  - Because all the nodes in the graph have degree of 2
  - The computational graphs will be the same binary tree


$v_1$  resides in a cycle with length 3




$v_1$  resides in a cycle with length 4



There exists no cycle that includes  $v_1$



The computational graphs for node  $v_1$  are always the same



# Feature Augmentation on Graphs (5)

- **Solution:**

- We can use **cycle count** as augmented node features


We start from cycle with length 0

Augmented node feature for  $v_1$

**[0, 0, 0, 1, 0, 0]**



$v_1$  resides in **a cycle with length 3**




Augmented node feature for  $v_1$

**[0, 0, 0, 0, 1, 0]**



$v_1$  resides in **a cycle with length 4**




# Feature Augmentation on Graphs (6)

- Other commonly used augmented features:
  - **Clustering coefficient**
  - **PageRank**
  - **Centrality**
  - ...
- **Any feature we have introduced in Lecture 2 can be used!**

# Add Virtual Nodes / Edges (1)


- **Motivation:** Augment sparse graphs
- **(1) Add virtual edges**
  - **Common approach:** Connect 2-hop neighbors via virtual edges
  - **Intuition:** Instead of using adj. matrix  $A$  for GNN computation, use  $A + A^2$
  - **Use cases:** Bipartite graphs
    - Author-to-papers (they authored)
    - 2-hop virtual edges make an author-author collaboration graph





# Add Virtual Nodes / Edges (2)

- **Motivation:** Augment sparse graphs
- **(2) Add virtual nodes**
  - The virtual node will connect to all the nodes in the graph
    - Suppose in a sparse graph, two nodes have shortest path distance of 10
    - After adding the virtual node, **all the nodes will have a distance of 2**
      - Node A – Virtual node – Node B
  - **Benefits:** Greatly **improves message passing in sparse graphs**



# Node Neighborhood Sampling

- **Previously:**

- All the nodes are used for message passing



- **New idea:** (Randomly) sample a node's neighborhood for message passing

# Neighborhood Sampling Example (1)

- For example, we can randomly choose 2 neighbors to pass messages
  - Only nodes *B* and *D* will pass message to *A*



# Neighborhood Sampling Example (3)

- In expectation, we can get embeddings similar to the case where all the neighbors are used
  - **Benefits:** greatly reduce computational cost
  - And in practice it works great!
  - We will talk more about scalable GNNs in Lecture 6



# Summary of the Lecture

- **Recap: A general perspective for GNNs**
  - **GNN Layer:**
    - Transformation + Aggregation
    - Classic GNN layers: GCN, GraphSAGE, GAT
  - **Layer connectivity:**
    - Deciding number of layers
    - Skip connections
  - **Graph Manipulation:**
    - Feature augmentation
    - Structure manipulation
- Having understood the basics of GNNs, we can now explore advanced architectures, tasks and applications!