

# Trustworthy AI for Graphs

CPSC483: Deep Learning on Graph-Structured Data

Rex Ying

# Readings

- Readings are updated on the website (syllabus page)
- **PyG Lecture Readings:**
  - [Documentation](#)
  - [GitHub](#)
- **Lecture 15 readings:**
  - [Trustworthy Graph Neural Networks](#)
  - [GraphFramEx](#) Evaluation

# Outline of Today's Lecture

- 1. Intro to Trustworthy GNN**
- 2. Adversarial Attacks and Robustness of GNNs**
- 3. Explainability for GNNs**

# Outline of Today's Lecture

## **1. Intro to Trustworthy GNN**

## 2. Adversarial Attacks and Robustness of GNNs

## 3. Explainability for GNNs

# Trustworthy Graph Learning

- **Trustworthy** AI/GNN includes many components
  - Explainability, fairness, robustness, privacy, ...
  - Algorithms to tackle combination of these aspects
- **Challenges**
  - Role of graph topology is previously unexplored in these problems.
  - Comprehensive quantitative evaluation

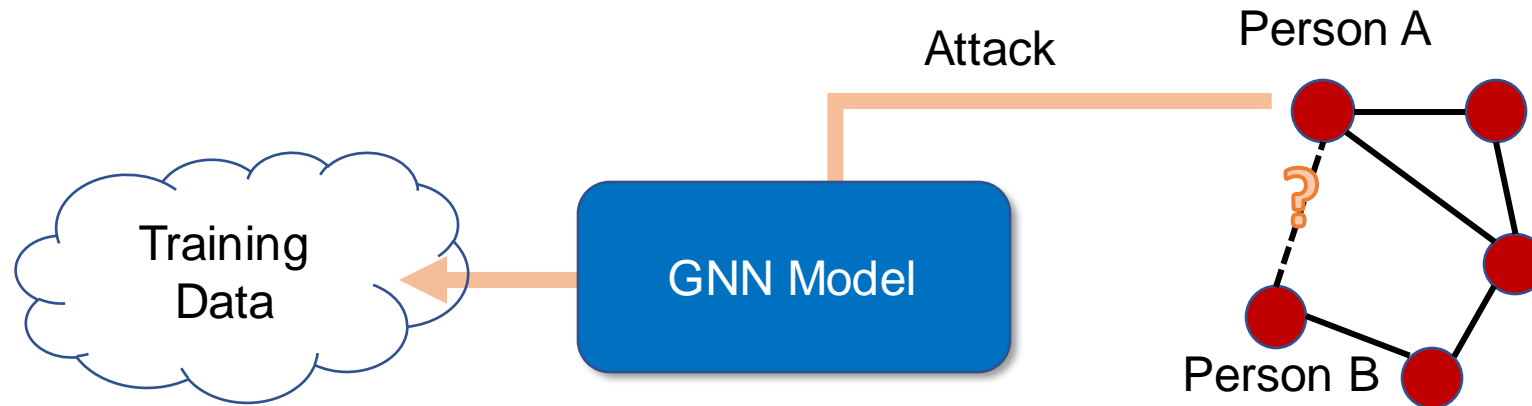
# Aspects of Trustworthy GNNs

- **Robustness** (part 2)
- **Explainability** (part 3)
- Privacy
- Fairness
- Environmental well-being
- Others

**How does each aspect play a role in gaining trust from users of machine learning models?**

# Privacy of GNNs

- Prevent private data within GNNs from being leaked
  - Training data
  - Model parameters
- Ex) Query social relations between two people by attacking the GNN model



# Privacy of GNNs: Attack

- **Model extraction attacks**
  - Steal architecture and parameters of a GNN model.
- **Membership inference attacks**
  - Infer whether a node/link/graph belongs to the training set of a GNN model.
- **Model inversion attacks**
  - Infer a GNN model's inputs from their corresponding outputs.
- **Other privacy attacks**



# Privacy of GNNs: Privacy-Preserving Techniques

- **Federated Learning**

- Calculate gradients on **individuals** using their own data
- Aggregate parameters (e.g. gradients/model weights) on the server

- **Differential Privacy**

- Add noise to data, such that
  - **Meaningless** when viewed individually
  - But approximate the analytics results when **aggregated**
- Noise can cancel out when performing the mean aggregation operation in GNN layers.

- **Insusceptible Training**

Original task loss    Attack function: try to distinguish the private labels

$$\min_{\theta} \sum_{v_i \in \mathcal{V}} \mathcal{L}_Y(f_{\theta}(v_i)) + \lambda \mathcal{L}_A(\mathcal{F}_A(v_i))$$

- **Security Computation**

- More related to system/hardware

Privacy-preserving loss: e.g. make the attack function's output probability close to 0.5 for the private labels

# Fairness of GNNs

- **Goal:** exclude prejudice or favoritism towards an individual or a group.
  - For example, in a bank's transaction network, the model should not learn to make predictions of loans based on gender, race or other protected characteristics.
- **Prevent Bias & Discrimination**
  - **Bias:** unfair operation in data collection, sampling, measurements, ...
  - **Discrimination:** incorporation of intentional or unintentional human prejudices and stereotyping in deep learning models

# Fairness of GNNs: Methods

- **Fair representation learning methods**

- Learn representations, from which one cannot infer sensitive attributes.
- A common technique is **adversarial training**

- **Fair prediction enhancement methods**

- **Data augmentation**

- Perturbation of protected features

- **Fair graph**

- Modify graph structures (e.g. drop edges that may induce bias)

- **Regularisation**

- Ex) any two individuals who are similar should receive similar algorithmic outcome

$$\| \mathbf{Y}[i, :] - \mathbf{Y}[j, :] \|_F^2 \mathbf{S}[i, j] \leq \delta$$

Predictions of node  $i$       Similarity between node  $i$  and  $j$

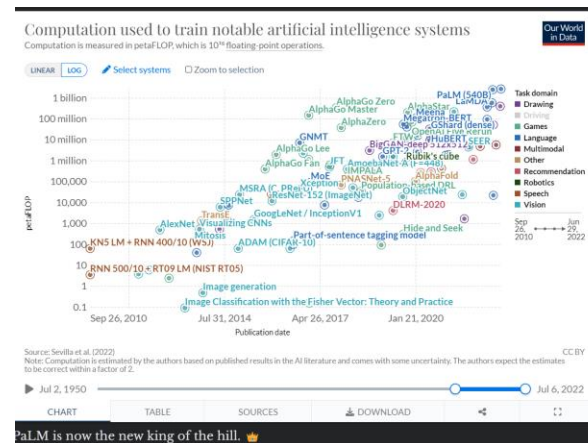
# Environmental Well-Being of GNNs

- GNNs should conform to **the fundamental values of the society** in which they are deployed
  - Large-scale graph datasets → GNN execution efficiency
  - Deeper or more complex architectures → GNN's deployment on edge devices
  - The unique characteristics of graph data → specially designed software and hardware



**Compute Clusters**

Credit: Imaginima/E+/gettyimages



**Large model training involves  $10^{24}$  flops**

<https://blog.heim.xyz/palm-training-cost/>

tl;dr What would it cost you to train PaLM using cloud computing (and you're not Google)?  
Something around **\$9M to \$23M**.

## PaLM a 540B state-of-the-art language model

Google recently published a new paper presenting PaLM (their blogpost) – a 540B parameter large language model.

**Input:** Jennifer looked out her window and sees a really cool cloud below her. She unbuckles her seatbelt and heads to the bathroom. Is Jennifer probably traveling more than 300 miles per hour relative to the earth?  
**Model Output:** 300 miles per hour is about 480 km/h. This is about the speed of a commercial airplane. Clouds are usually below airplanes, so Jennifer is probably on an airplane. The answer is "yes".

**540 B parameter pretrained model**

<https://blog.heim.xyz/palm-training-cost/>

# Environmental Well-Being of GNNs: Methods

- **Scalable GNN**
  - We've already talked about it in [Lecture 6](#)
- **Model compression**
  - Knowledge distillation
  - Model pruning
  - Reducing parameters
  - Model quantisation
- **Efficient frameworks and accelerators**
  - Software: [PyTorch Geometric](#), **Efficient AutoML**
  - Hardware: [EnGN](#), [HyGCN](#)

# Outline of Today's Lecture

1. Trustworthy AI

**2. Adversarial Attacks and Robustness of GNNs**

3. Explainability for GNNs

# Deep Learning Performance

- Recent years have seen **impressive performance of deep learning models in a variety of applications.**
  - In computer vision, **deep convolutional networks** have achieved human-level performance on ImageNet (image category classification)
- **Are these models ready to be deployed in real world?**

# Adversarial Examples

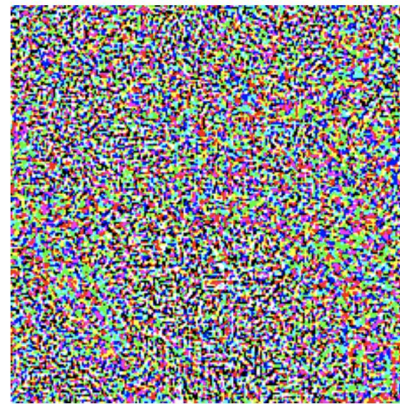
- Deep convolutional neural networks are vulnerable to **adversarial attacks**:
  - Imperceptible noise changes the prediction.



“Panda”

57.7% confidence

+ .007 ×



Carefully

calculated noise

=



“Gibbon”

93.3% confidence

Goodfellow, I., Shlens, J., & Szegedy, C.. (2014). Explaining and Harnessing Adversarial Examples.

- Adversarial examples are also reported in natural language processing [Jia & Liang et al. EMNLP 2017] and audio processing [Carlini et al. 2018] domains.



# Implication of Adversarial Examples

- **The existence of adversarial examples prevents the reliable deployment of deep learning models to the real world.**
  - Adversaries may try to actively hack the deep learning models.
  - The model performance can become much worse than we expect.
- **Deep learning models are often not robust.**
  - It is an active area of research to make these models robust against adversarial examples

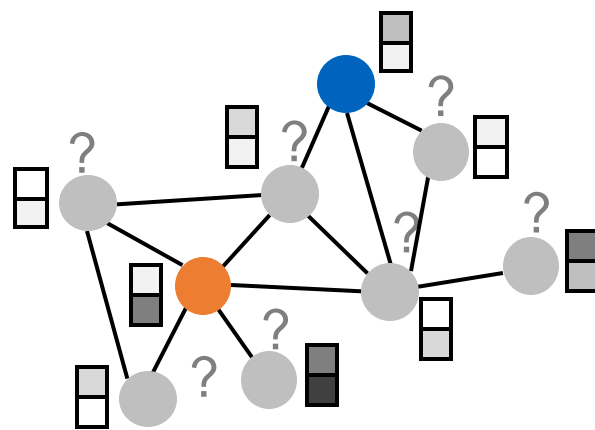
# Robustness of GNNs

- **This lecture: How about GNNs? Are they robust to adversarial examples?**
- **Premise:** Common applications of GNNs involve **public platforms** and **monetary interests**.
  - Recommender systems
  - Social networks
  - E-commerce platforms
- **Adversaries have the incentive to** manipulate input graphs and hack GNNs' predictions.

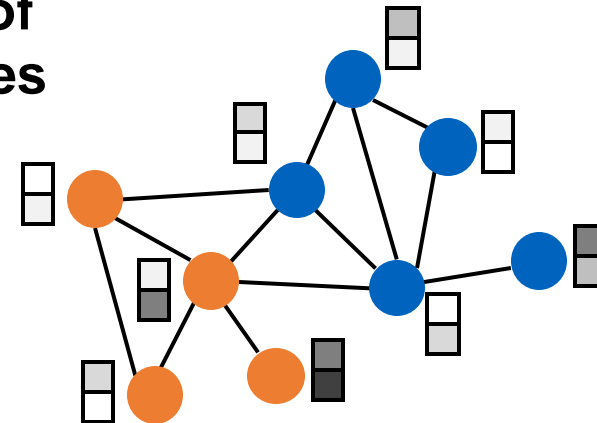
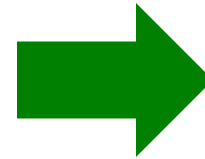
# Setting to Study GNNs' Robustness

- To study the robustness of GNNs, we specifically consider the following setting:
  - **Task**: Semi-supervised node classification
  - **Model**: GCN [Kipf & Welling ICLR 2017]

?: Unlabeled



**Predict labels of unlabeled nodes**

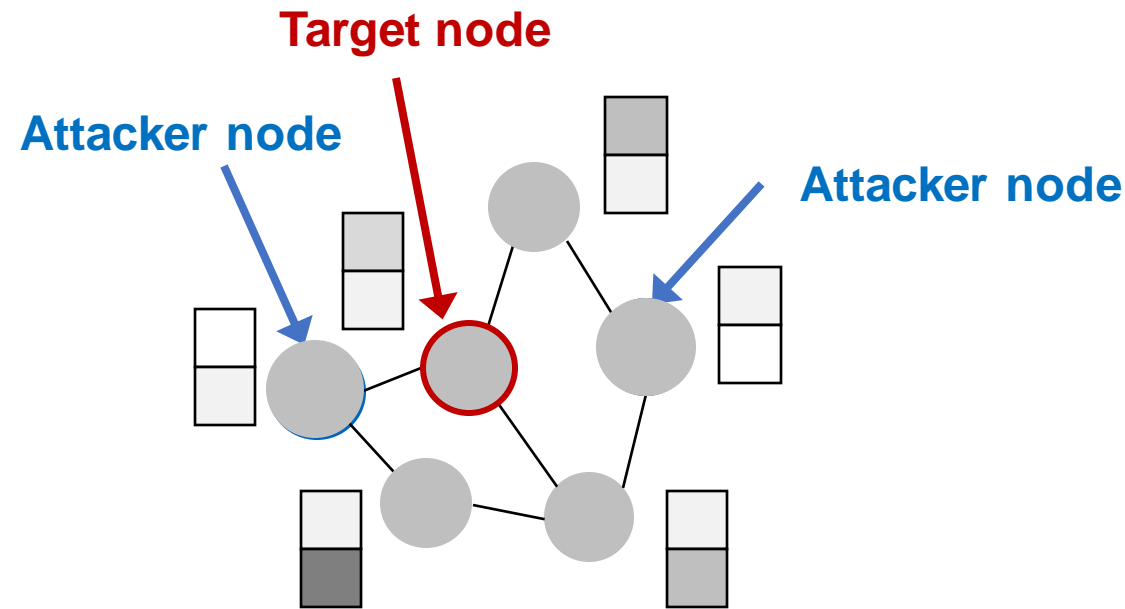


# Roadmap

- We first describe several real-world **adversarial attack possibilities**.
- We then review the GCN model that we are going to attack (**knowing the opponent**).
- We mathematically **formalize the attack problem as an optimization problem**.
- **We empirically see how vulnerable GCN's prediction is to the adversarial attack.**

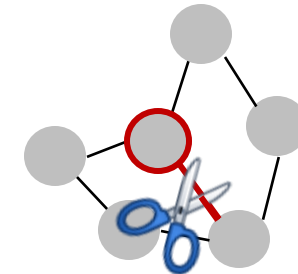
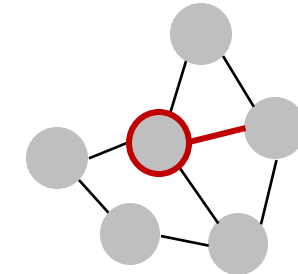
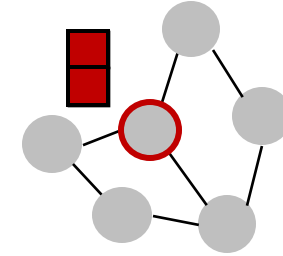
# Attack Possibilities

- What are the attack possibilities in real world?
  - **Target node**  $t \in V$ : node whose label prediction we want to change
  - **Attacker nodes**  $S \subset V$ : nodes the attacker can modify



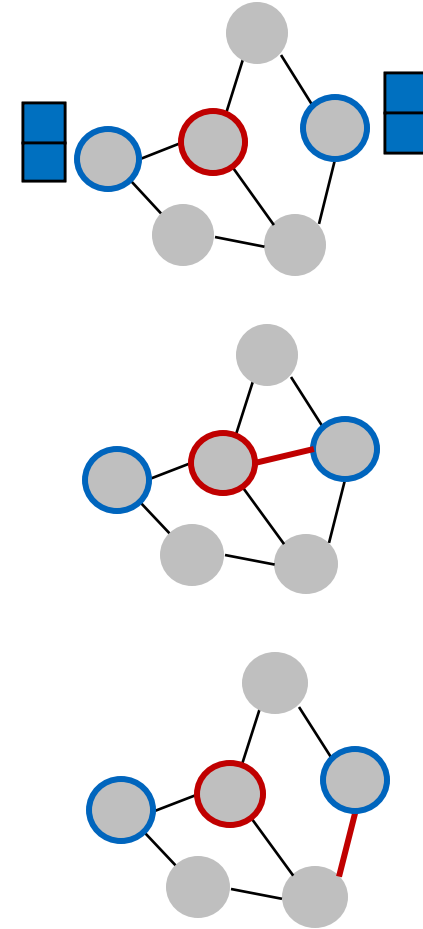
# Attack Possibilities: Direct Attack

- **Direct Attack: Attacker** node is the **target** node:  $S = \{t\}$
- Modify **target** node feature
  - Change website content
- Add connections to **target**
  - Buy/likes/follows
- Remove connections from **target**
  - Unfollow users



# Attack Possibilities: Indirect Attack

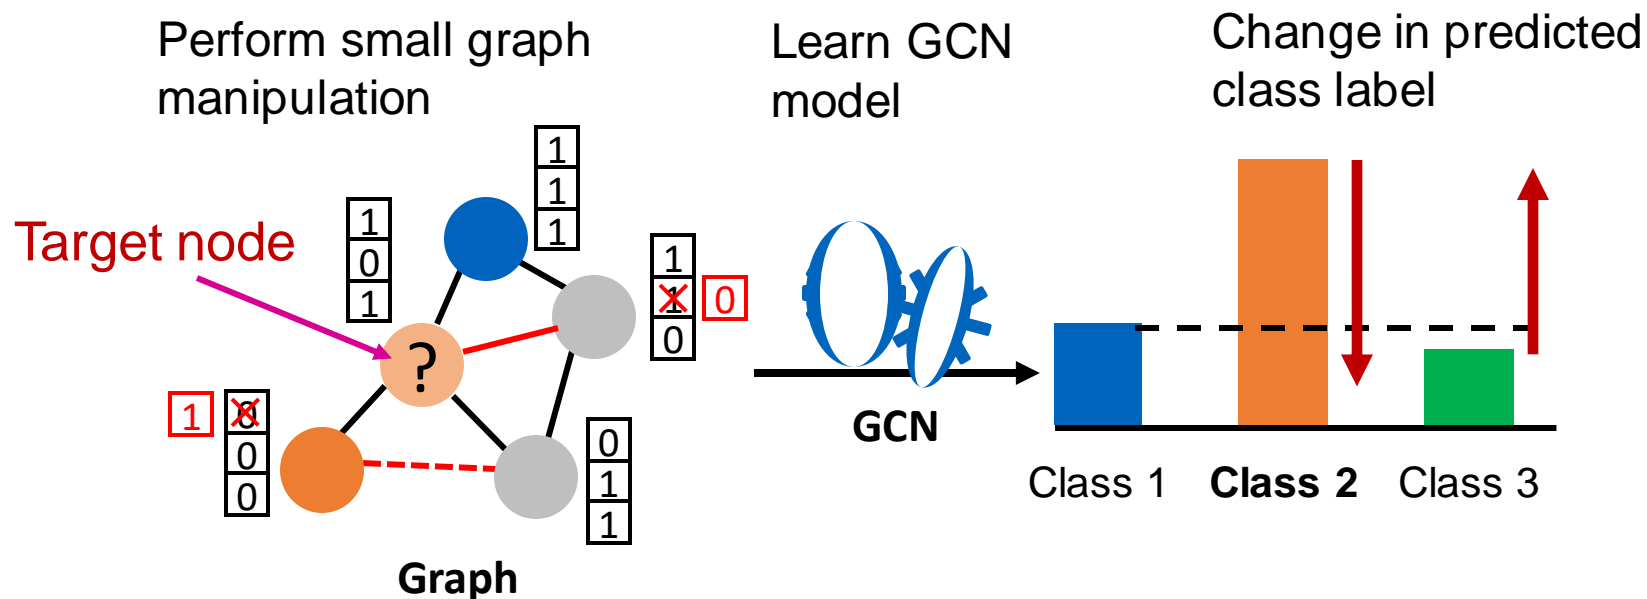
- **Indirect Attack:** The **target** node is not in the **attacker** nodes:  $t \notin S$
- Modify **attacker** node features
  - Ex) Hijack friends of targets
- Add connections to **attackers**
  - Ex) Create a link, link farm
- Remove connections from **attackers**
  - Ex) Delete undesirable link



# Formalizing Adversarial Attackers

- Objective for the attacker:

Maximize (**change of target node label prediction**)  
Subject to (**graph manipulation is small**)



If graph manipulation is too large, it will easily be detected. Successful attacks should change the target prediction with “unnoticeably-small” graph manipulation.



# Mathematical Formulation (1)

- **Original graph:**
  - $A$ : adjacency matrix,  $X$ : feature matrix
- **Manipulated graph (after adding noise):**
  - $A'$ : adjacency matrix,  $X'$ : feature matrix
- **Assumption:**  $(A', X') \approx (A, X)$ 
  - Graph manipulation is **unnoticeably small**.
    - Preserving basic graph statistics (e.g., degree distribution) and feature statistics.
  - Graph manipulation is either **direct** (changing the feature/connection of target nodes) or **indirect**.

# Mathematical Formulation (2)

- **Target node:**  $v \in V$

- Recall that we only consider semi-supervised node classification settings

- GCN learned over the **graph**

$$\theta^* = \operatorname{argmin}_{\theta} \mathcal{L}_{train}(\theta; \tilde{A}, \tilde{X})$$

What is  $\tilde{A}$  and  $\tilde{X}$  ?

- GCN's original prediction on the **target node**:

$$c_v^* = \operatorname{argmax}_c f_{\theta^*}(A, X)_{v,c}$$

Predict the class  $c_v^*$  of vertex  $v$  that has the highest predicted probability

- GCN's manipulated prediction on the target node

$$c_v^{*'} = \operatorname{argmax}_c f_{\theta^*}(A', X')_{v,c}$$

- **We want the prediction to change after the graph is manipulated:**

$$c_v^{*'} \neq c_v^*$$

# Mathematical Formulation (3)

- What are  $\tilde{A}$ ,  $\tilde{X}$  ?
- Evasion Attack:  $\tilde{A} = A$ ,  $\tilde{X} = X$ 
  - Attacking happens after the GNN model is trained (at test time)
  - Inductive setting
- Poisoning Attack:  $\tilde{A} = A'$ ,  $\tilde{X} = X'$ 
  - Attacking happens before the GNN model is trained.

# Mathematical Formulation (4)

- Change of prediction on target node  $v$ :

$$\Delta(v; A', X') =$$

$$\log f_{\theta^*}(A', X')_{v, c_v^{*'}} - \log f_{\theta^*}(A', X')_{v, c_v^*}$$

Predicted (log)  
probability of the  
newly-predicted  
class  $c_v^{*'}$



Want to increase  
this term

Predicted (log)  
probability of the  
originally-predicted  
class  $c_v^*$



Want to decrease  
this term

# Mathematical Formulation (5)

- **Final optimization objective:**

$$\begin{aligned} & \operatorname{argmax}_{A', X'} \Delta(v; A', X') \\ & \text{subject to } (A', X') \approx (A, X) \end{aligned}$$

- **Challenges in optimizing the objective**

- Adjacency matrix  $A'$  is a discrete object: gradient-based optimization cannot be used.
- Several approximations are proposed to make the optimization.

# Nettack: Greedy Scheme

- A **greedy** scheme
- while  $|A' - A| + |X' - X| < \delta$ 
  - Compute a **candidate set** for **struct perturbations** (all should be **unnoticeable**)
  - Pick the one which obtains the **highest** change of predictions
  - Compute a candidate set for **node features perturbations** (all should be **unnoticeable**)
  - Pick the one which obtains the **highest** change of predictions
- Challenges in this scheme:
  - How to make sure the perturbation is **unnoticeable** (budget  $\delta$  is not enough)?
  - How to efficiently compute the **candidate sets**?
  - How to efficiently get the one which make the **highest change**?

Manipulation budget

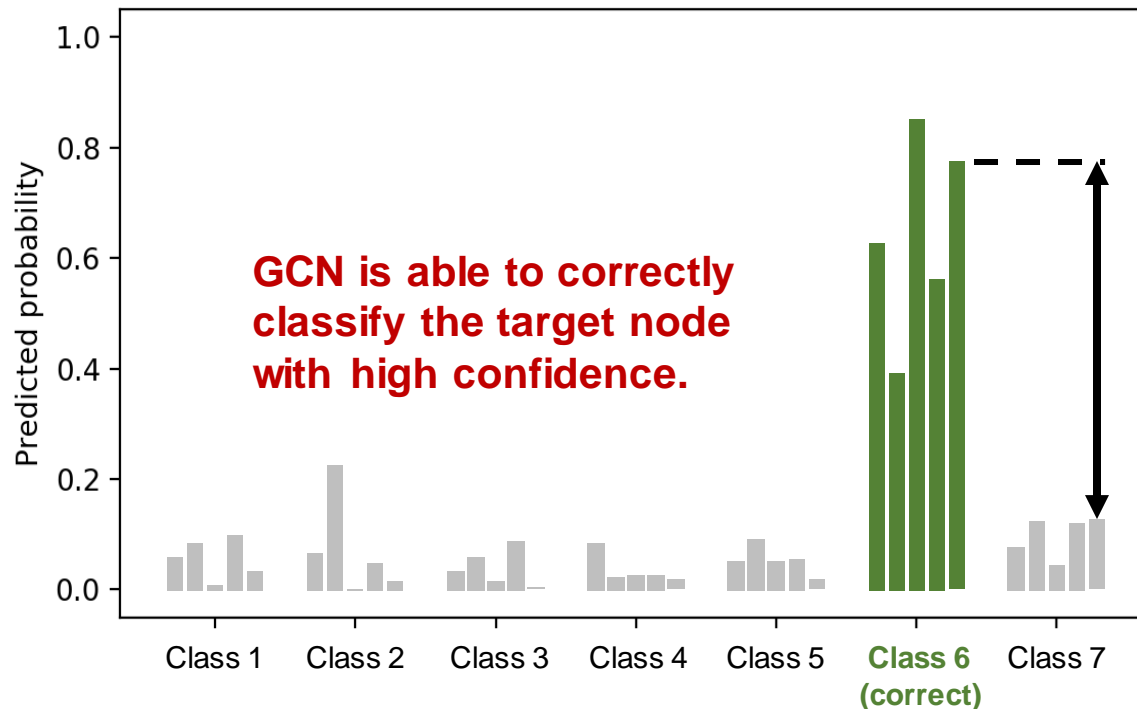
$\delta$

# Netattack: Candidate Sets

- Efficiently compute **candidate sets** which make **unnoticeable** changes
- Main idea: use **significance test**
- **Graph structure** preserving perturbations
  - Preserve a graph's degree distribution (usually follow  $p(x) \propto x^{-\alpha}$ )
  - Can be **incrementally computed (constant time)** during changes
- **Feature statistics** preserving perturbations
  - Preserve feature co-occurrence
  - Can be **precomputed**
- Find more details in [\[Zügner, KDD 2018\]](#)

# Experiments: Adversarial Attack (1)

- **Semi-supervised node classification with GCN on a paper citation network (2,800 nodes, 8,000 edges).**



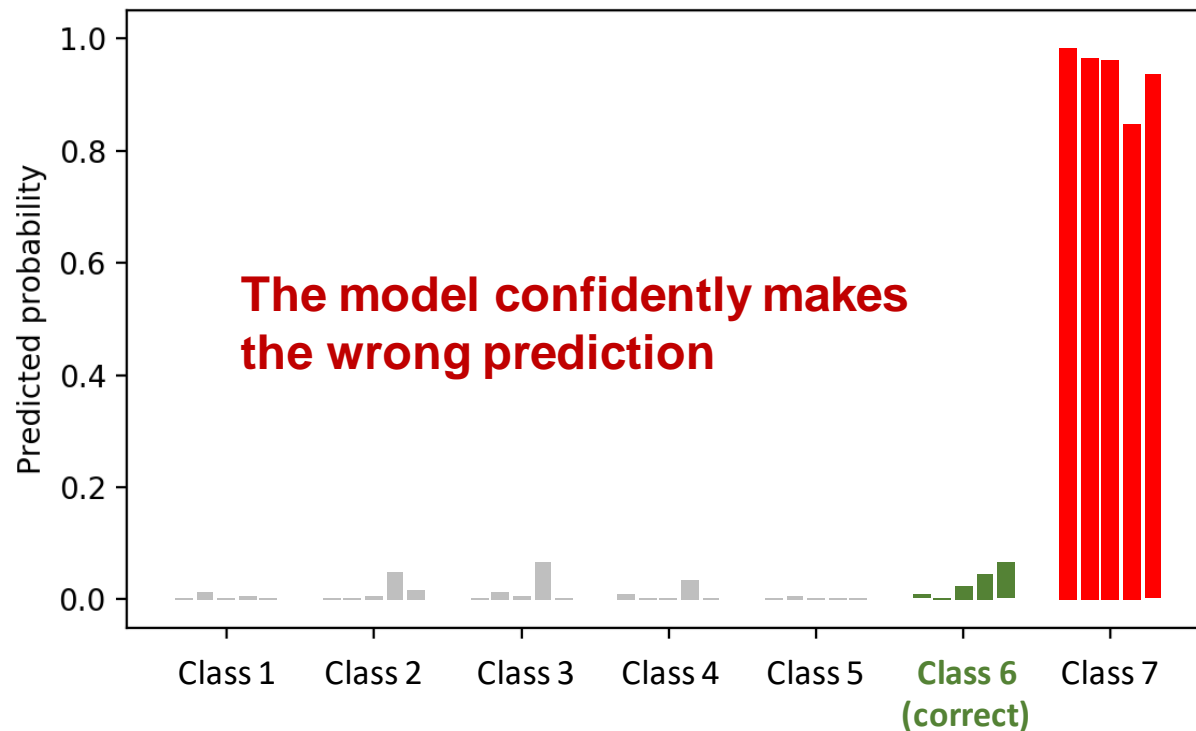
Predicted probabilities over 5 re-trainings  
(without graph manipulation, i.e., clean graph)

**Classification margin**  
> 0: Correct classification  
< 0: Incorrect classification



# Experiments: Adversarial Attack (2)

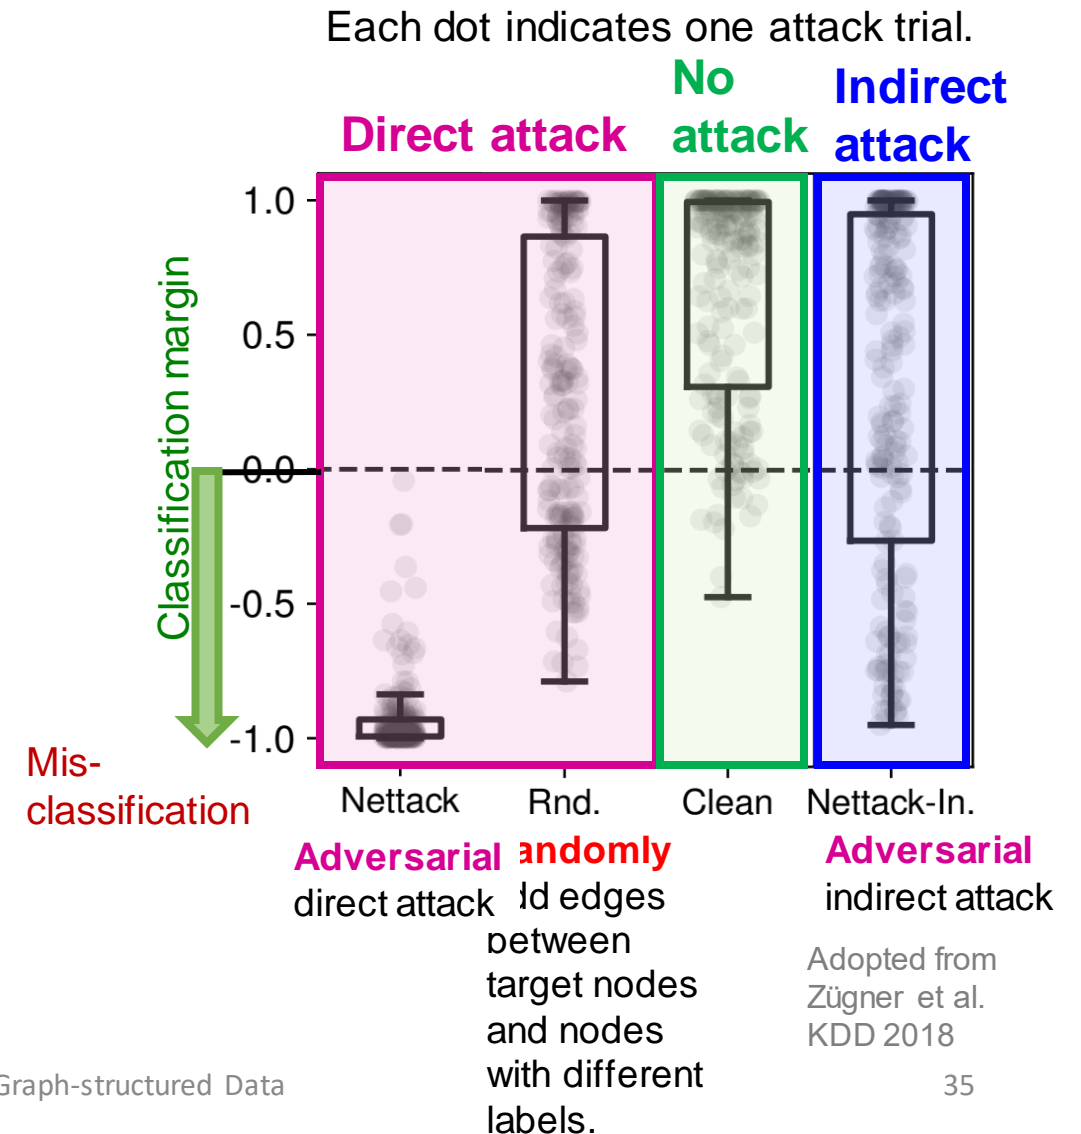
- GCN's prediction after carefully modifying just 5 edges attached to the target node (**direct adversarial attack**).



Predicted probabilities over 5 re-trainings  
(with adversarial attacks)

# Experiments: Attack Comparison

- **Adversarial direct attack** is the strongest attack, significantly worsening GCN's performance (compared to **no attack**).
- **Random** attack is much weaker than **adversarial** attack.
- **Indirect attack** is more challenging than direct attack.



# Outline of Today's Lecture

1. Trustworthy AI

2. Adversarial Attacks and Robustness of GNNs

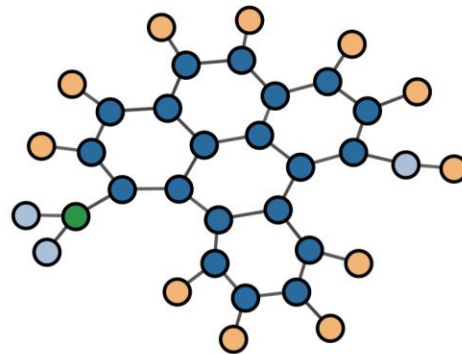
**3. Explainability for GNNs**

# Explainability: Motivation (1)

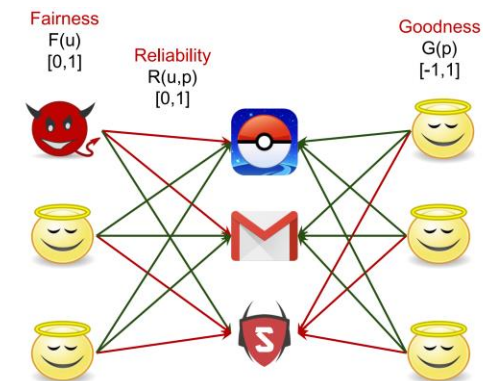
- Example questions after training GNNs:
  - Why is an item recommended to a user?
  - Why is the molecule mutagenic?
  - Why is the user classified as fraudulent?
- How to let the domain experts understand and trust the GNN model?



Recommender System



Mutagenic Molecule



Fraudulent Detection

# Explainability: Motivation (2)

- **Example questions after training GNNs:**
  - Why is an item recommended to a user?  
Explain link prediction
  - Why is the molecule mutagenic?  
Explain graph classification
  - Why is the user classified as fraudulent?  
Explain node classification
- **Need to provide explanations to GNN models!**

# Deep Learning Explainability Methods: Examples

- **Proxy Model**

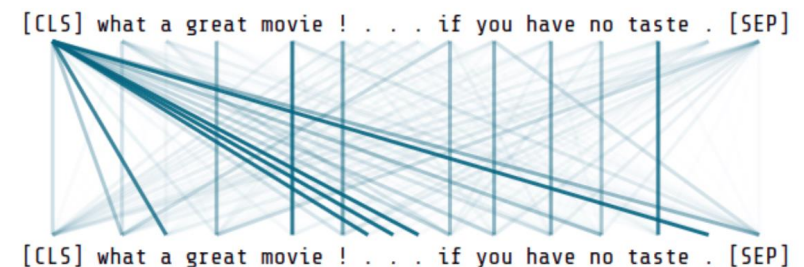
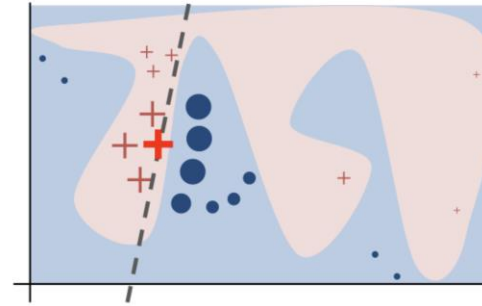
- Learn a inherently interpretable model locally approximating the original model (e.g. a linear model, interpret by weights).

- **Saliency Maps**

- Compute gradients of outputs w.r.t. input pixels.

- **Attention Mechanisms**

- Visualize attention weights in a attention models.

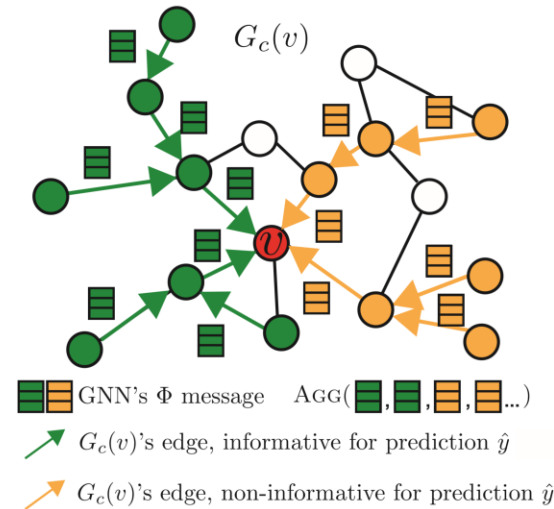


# Challenges of Applying these Methods to Graphs

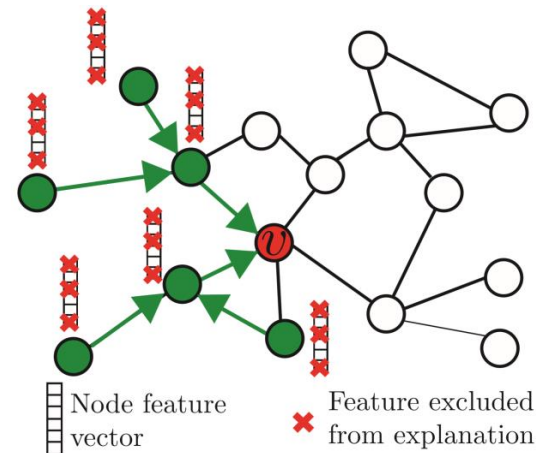
- Why is it non-trivial to apply these explainability methods to graph structure?
  - predictions on graphs are induced by a complex combination of nodes and paths of edges between them (not only nodes).

# How to explain a GNN

- Message passing structure
- The importance of node features



**Structural explanation**



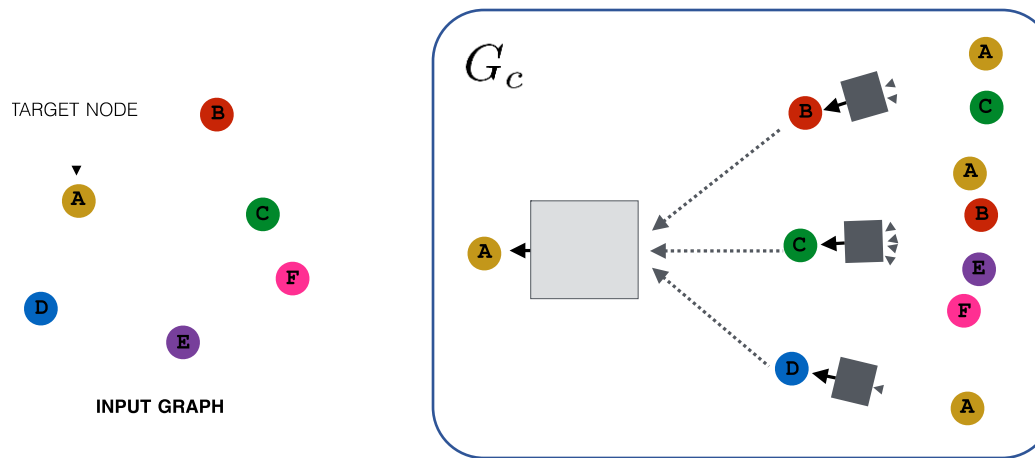
**Feature explanation**

- GNNExplainer explain both aspects **simultaneously**



# GNNExplainer Input

- Without loss of generality, consider node classification task:

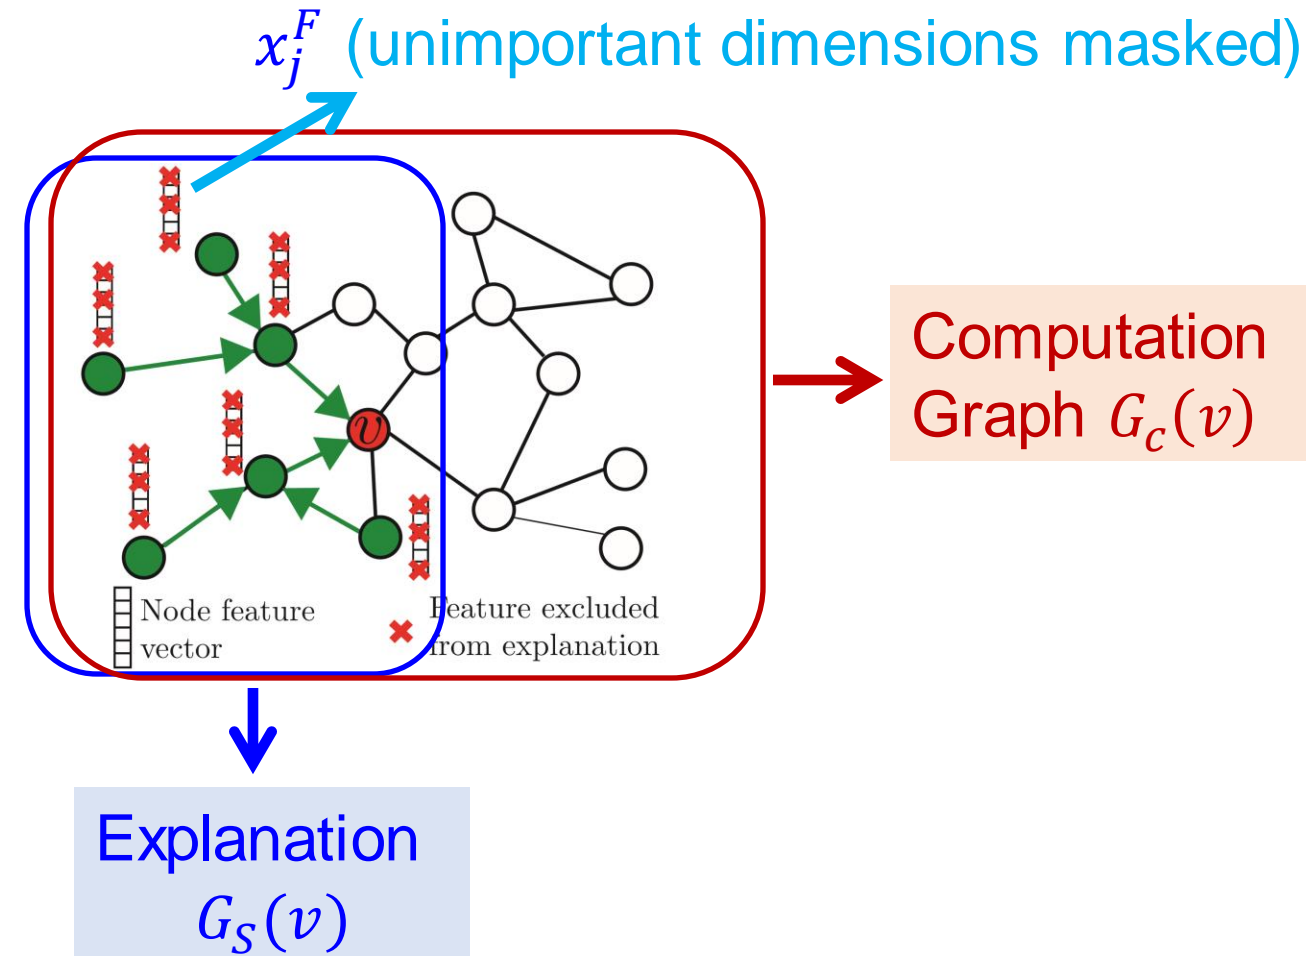


Suppose GNN predicts  
label  $\hat{y}$  for node  $v$

- Input computation graph:  $G_c(v)$
- Adjacency matrix of  $G_c$ :  $A_c(v) \in \{0,1\}^{n \times n}$
- Node Feature:  $X_c(v) = \{x_j | v_j \in G_c(v)\}$

# GNNExplainer Output

- GNN model  $\phi$  learns  $P_{\phi}(Y \mid A_c(v), X_c(v))$
- $Y$  denotes predicted label of  $v$
- **GNNExplainer** outputs  $(A_S, X_S^F)$
- Graph  $G_S$  with adjacency matrix  $A_S$  is a subgraph of graph with adjacency matrix  $A_c(v)$  (omit  $v$ )
- $X_S^F = \{x_j^F \mid v_j \in G_S\}$  are features for  $G_S$
- Mask  $F$  masks out unimportant dimensions



# Explain by Mutual Information

- **Mutual information (MI)**

- A measure of the mutual correlation between the two random variables.
- Good explanation should have **high correlation** with model prediction
- **Relation to entropy:**

$$MI(X; Y) = H(X) - H(X|Y) = H(Y) - H(Y|X)$$

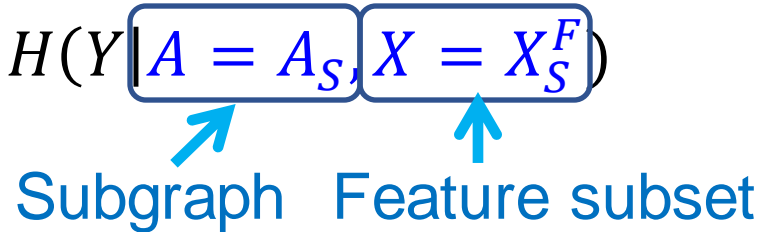
- GNNExplainer **Objective:**

- **Maximize MI** between **label** and **explanation**

$$\max_{G_S} MI(\mathbf{Y}; (\mathbf{A}_S, \mathbf{X}_S)) = H(Y) - H(\mathbf{Y} | \mathbf{A} = \mathbf{A}_S, \mathbf{X} = \mathbf{X}_S^F)$$

# Explain by Optimization

- By **relation to entropy**:  
equivalent to minimize conditional entropy

$$\max_{A_S} MI(Y|(A_S, X_S)) = \min_{A_S} H(Y|A = A_S, X = X_S^F)$$


- Finding  $A_S$  that minimizes the conditional entropy is **computationally expensive!**
  - Exponentially many possible  $A_S$
- **Solution**: Treat explanation as a distribution of “**plausible explanations**”, instead of a single graph
  - Optimize the expected explanation
  - **Benefit 1**: captures multiple possible explanations for the same node
  - **Benefit 2**: turns discrete optimization to continuous

# GNNExplainer Model

- continuous relaxation

- Optimize the expected adjacency matrix  $A_S$   
expectation of explanations

$$\min_{\mathcal{A}} \mathbb{E}_{A_S \sim \mathcal{A}} H(Y|A = A_S, X = X_S)$$

- View  $\mathbb{E}_{A_S \sim \mathcal{A}}$  as an adjacency matrix where entries are continuous

- Approximation

$$\min_{\mathcal{A}} H(Y|A = \mathbb{E}_{\mathcal{A}}[A_S], X = X_S)$$

- Optimize the expectation by masking

Element-wise multiply

- Use  $A_C \odot \text{Mask}$  to represent  $\mathbb{E}_{\mathcal{A}}[A_S]$

- If  $\text{Mask}_{ij}$  close to 1, keep edge  $(i, j)$ ; if close to 0, drop edge  $(i, j)$ .

# GNNExplainer Model

- Let  $\text{Mask} = \sigma(M)$  be the adjacency mask
  - Continuous relaxation:  $\sigma(M) \in \mathbb{R}$  instead of binary
  - **Sigmoid** function  $\sigma$  squashes  $M$  into  $[0, 1]$
  - Masking: Element-wise multiply  $\sigma(M)$  by  $A_c$
- Assume edges are independent

$$P_{\mathcal{A}}(A_S) = \prod_{e=(j,k) \in G_c(v_i)} A_S[j, k]$$

- Objective:

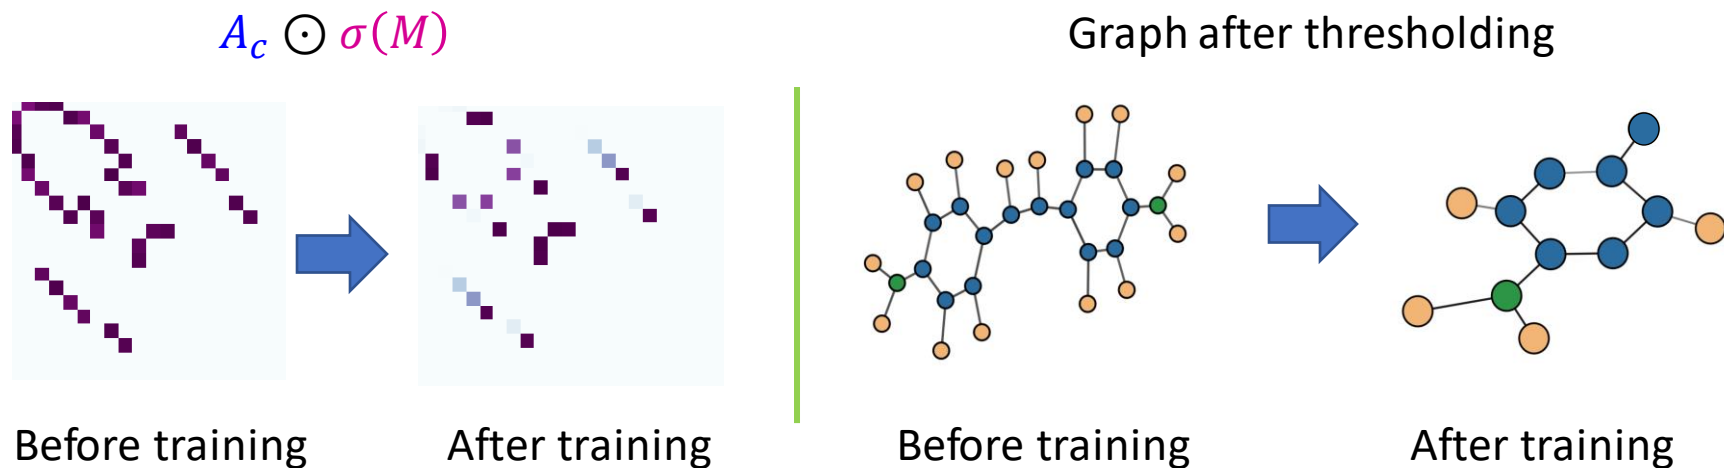
$$\min_M -H(\log P_{\phi}(Y = y | G = A_c \odot \sigma(M), X = X_S))$$

# GNNExplainer Model

- Optimize  $M$  :

$$\min_M -H(\log P_{\phi}(Y = y | A = A_c \odot \sigma(M), X = X_S))$$

- $A_c \odot \sigma(M)$  is the relaxed adjacency matrix
- Threshold  $A_c \odot \sigma(M)$  to get  $G_S$
- Example:



# Feature Explanation

- Similarly select features by **masking**

$$X_S^F = \{x_j^F \mid v_j \in G_S\}, \quad x_j^F = [x_{j,t_1}, \dots, x_{j,t_k}]$$

For the **selected dimensions**,  $\sigma(F_{t_i}) \rightarrow 1$

- **Problem**: Zero value could be important!
- **Solution**: Measure feature importance by how much drop in model confidence when features are replaced with random values.  
(See paper for details)



# Regularization Constraints

- Optimize feature and adjacency masks jointly with regularization

- **Concise explanation**

- Mask size:  $\text{Sum}(\sigma(M))$
- Feature size:  $\text{Sum}(\sigma(F))$

- **Final Objective**

$$\min_M -H(\log P_{\phi}(Y = y | G = A_c \odot \sigma(M), X = X_S^F)) + \lambda_1 \underbrace{\text{Sum}(\sigma(M)) + \lambda_2 \text{Sum}(\sigma(F))}_{\text{Sum of entries in feature and adjacency masks}}$$

- Threshold  $A_c \odot \sigma(M)$  to get the explanation  $G_S$

# GNNExplainer Model

- **Task extensions**

- Link prediction: optimize mask on union of 2 node neighborhoods
- Graph classification: optimize mask on graph

- **Can adapt to different architectures**

- Graph Attention Networks
- Gated Graph Sequence
- Graph Networks
- GraphSAGE
- Jumping Knowledge Networks

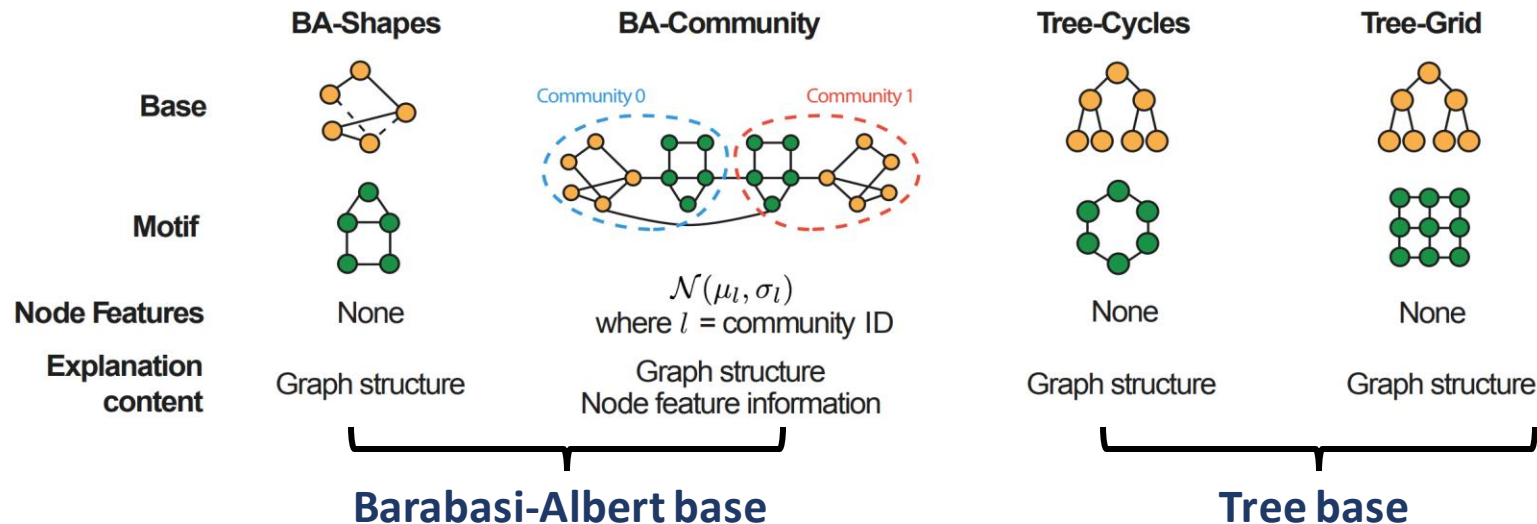
We use GNN  
general formulation

# Experiments: Baselines

- GNN saliency map based on gradients of output score w.r.t. inputs
- Attention values based on Graph Attention Networks (GAT)
  - Edge importance indicated by average attention weights across layers for each edge

# Experiments: Datasets (1)

- Synthetic task: **is a node part of a given motif?**
  - 100 Motifs are randomly attached to nodes in base graphs (500 nodes)
  - **Node classification (structural roles)**



# Experiments: Datasets (2)

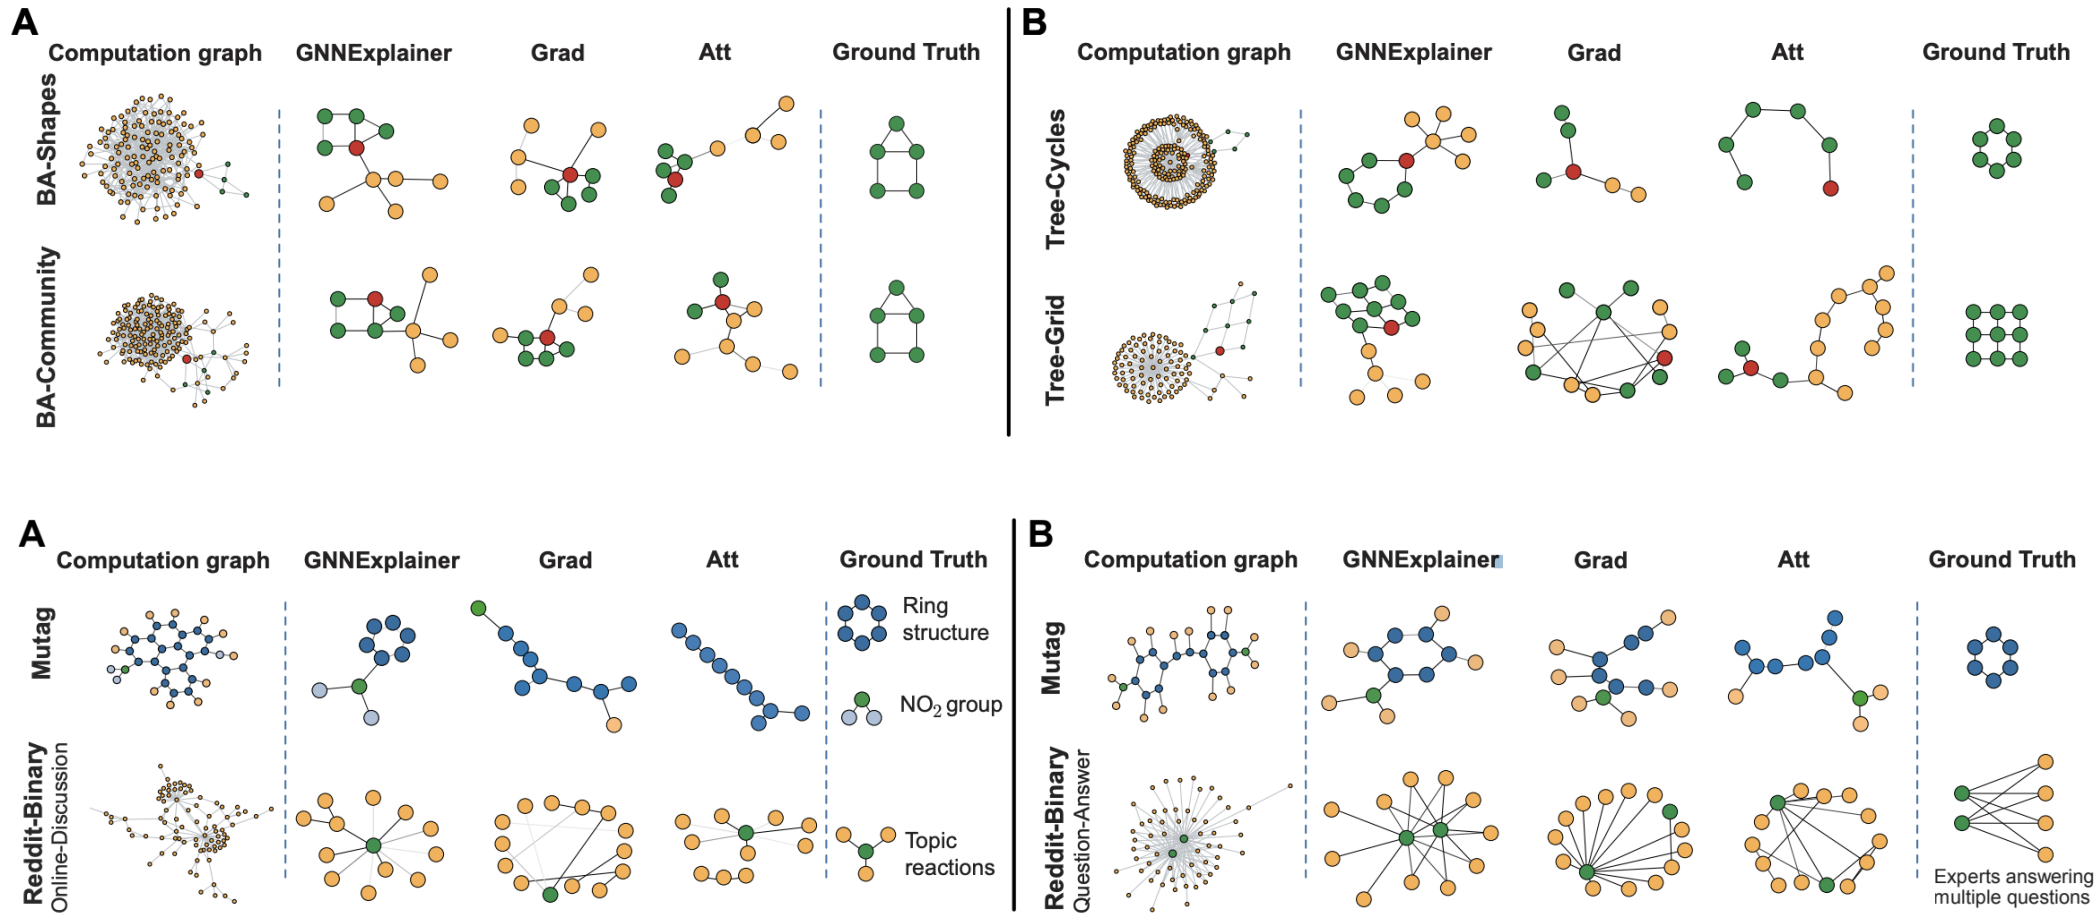
- Real-world tasks
  - Social networks (Reddit-binary dataset)
    - Reddit community prediction
  - Chemistry (Mutagenic molecule dataset)
    - Chemical property prediction
  - Graph classification

# Results: Quantitative Analysis

- Node classification with ground-truth
- Measures accuracy of **explanation** with respect to **ground-truth**

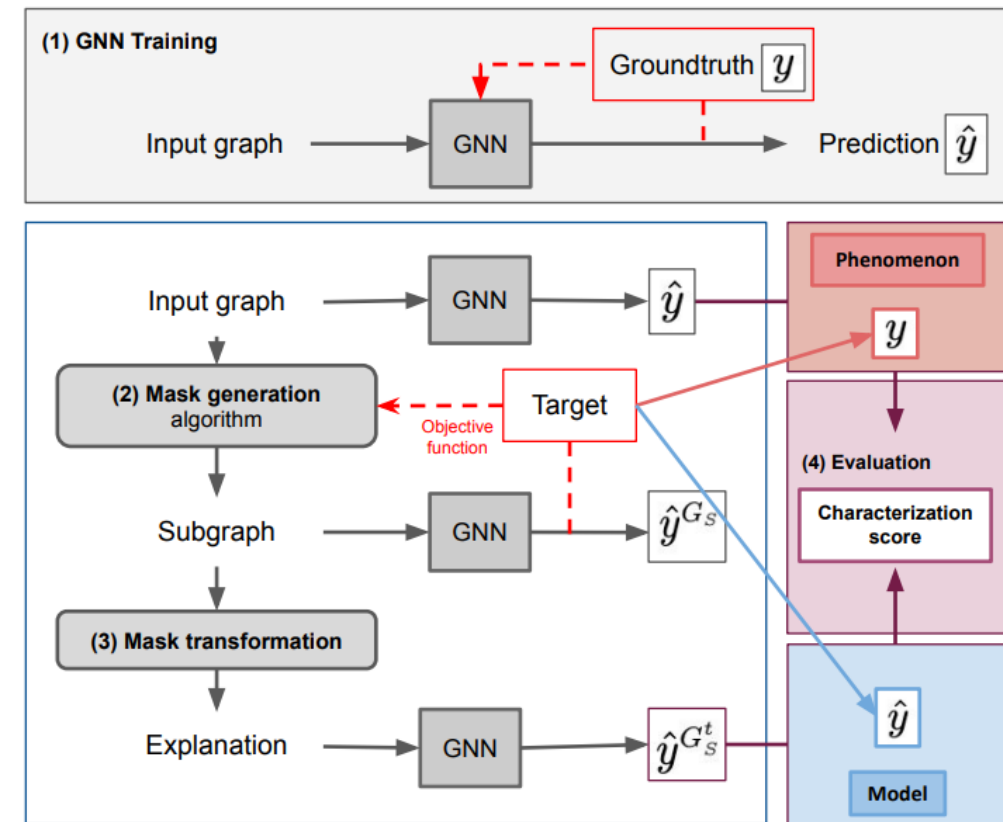
	BA-House	BA-Comm	Tree-Cycle	Tree-Grid
Grad	88.2	73.9	82.4	61.2
Att	81.5	75.0	90.5	66.7
<b>GNN-Explainer</b>	<b>92.5</b>	<b>83.6</b>	<b>94.8</b>	<b>87.5</b>

# Results: Qualitative Analysis



# Explainability Method Evaluation

- **Challenge:** **groundtruth** might not always be available
- Evaluation is **multi-dimensional**
- **Goal** (phenomenon vs. model)
- **Masking** strategy
- **Type** (sufficiency vs. necessity)
- **GraphFramEx**
- Benchmarks and evaluation criteria for graph explainability





# Explanation Goal

- **Phenomenon** Explanation
  - Explain the underlying reasons for the ground truth phenomenon
- **Model** Explanation
  - Explain why model makes a particular prediction
- Adapt the **fidelity** measure for both cases

Phenomenon

$$fid_+ = \frac{1}{N} \sum_{i=1}^N \left| \mathbb{1}(\hat{y}_i = y_i) - \mathbb{1}(\hat{y}_i^{G_{C \setminus S}} = y_i) \right|$$
$$fid_- = \frac{1}{N} \sum_{i=1}^N \left| \mathbb{1}(\hat{y}_i = y_i) - \mathbb{1}(\hat{y}_i^{G_S} = y_i) \right|$$

Model

$$fid_+ = 1 - \frac{1}{N} \sum_{i=1}^N \mathbb{1}(\hat{y}_i^{G_{C \setminus S}} = \hat{y}_i)$$
$$fid_- = 1 - \frac{1}{N} \sum_{i=1}^N \mathbb{1}(\hat{y}_i^{G_S} = \hat{y}_i)$$

# Masking Strategy

- **Hard mask**
  - Produce a subgraph as an explanation for the prediction
- **Soft mask**
  - A number between 0 to 1 on edges and features to indicate their importance

# Types of Explanations

- **Sufficiency**

- An explanation is sufficient if it leads by its own to the initial prediction of the model explanation. ( $fid_- \rightarrow 0$ )

- **Necessity**

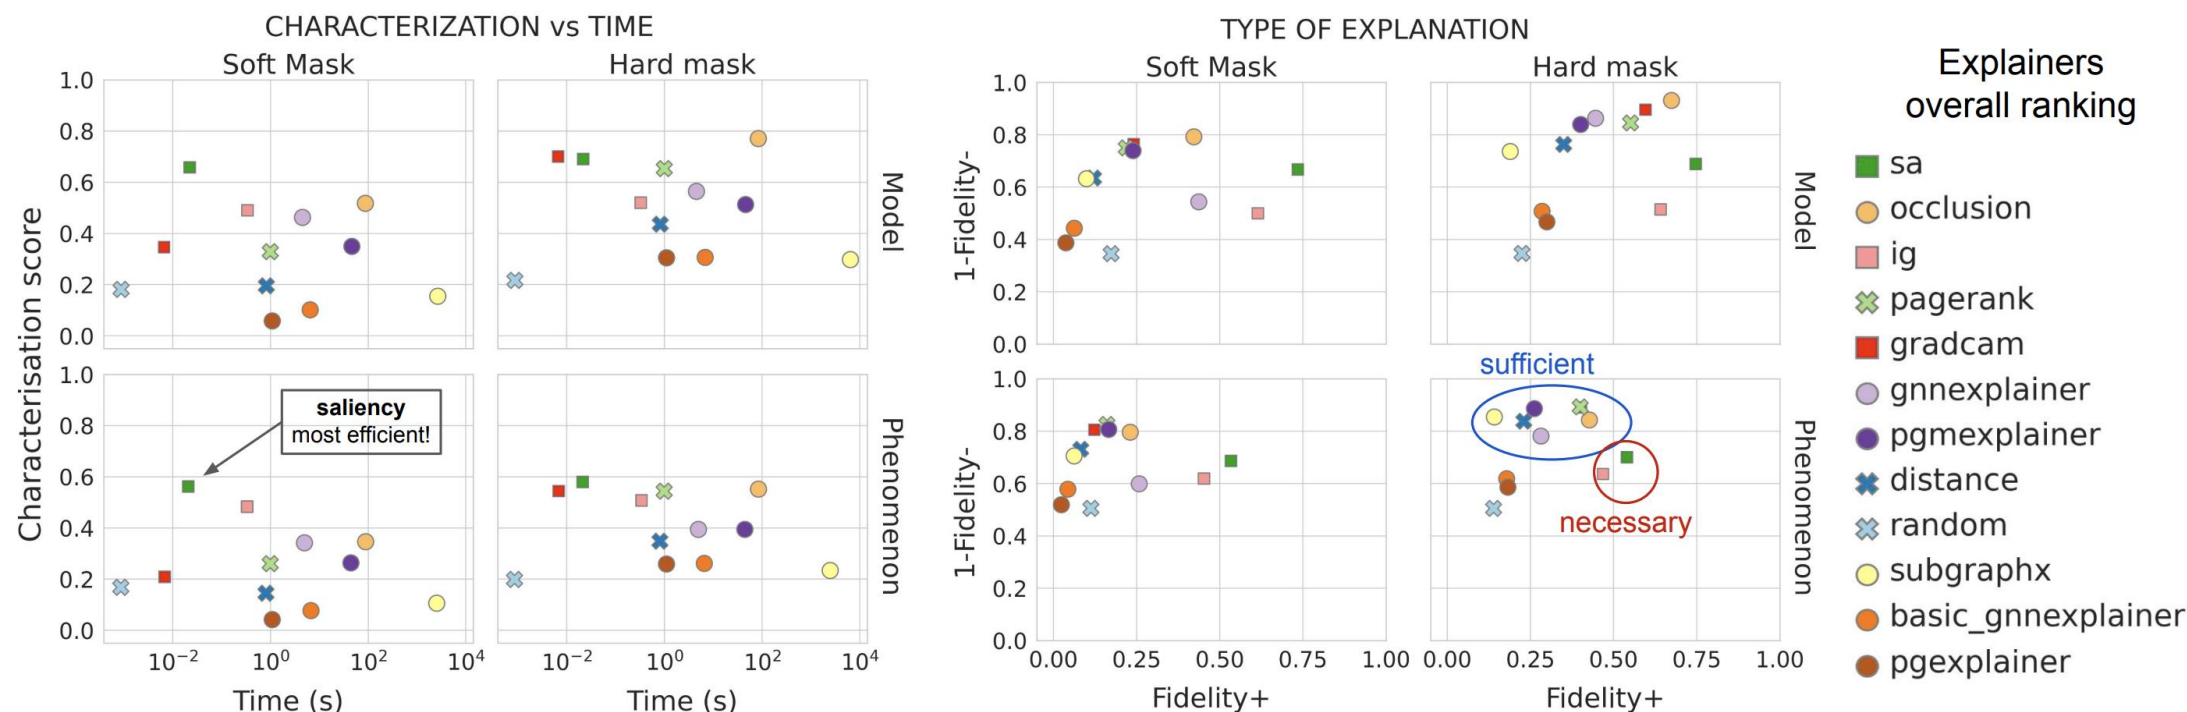
- An explanation is necessary if the model prediction changes when removing it from the initial graph. ( $fid_+ \rightarrow 1$ )

- **Characterization** score to summarize the explanation quality

$$character = \frac{w_+ + w_-}{\frac{w_+}{fid_+} + \frac{w_-}{1 - fid_-}} = \frac{(w_+ + w_-) \times fid_+ \times (1 - fid_-)}{w_+ \cdot (1 - fid_-) + w_- \cdot fid_+}$$

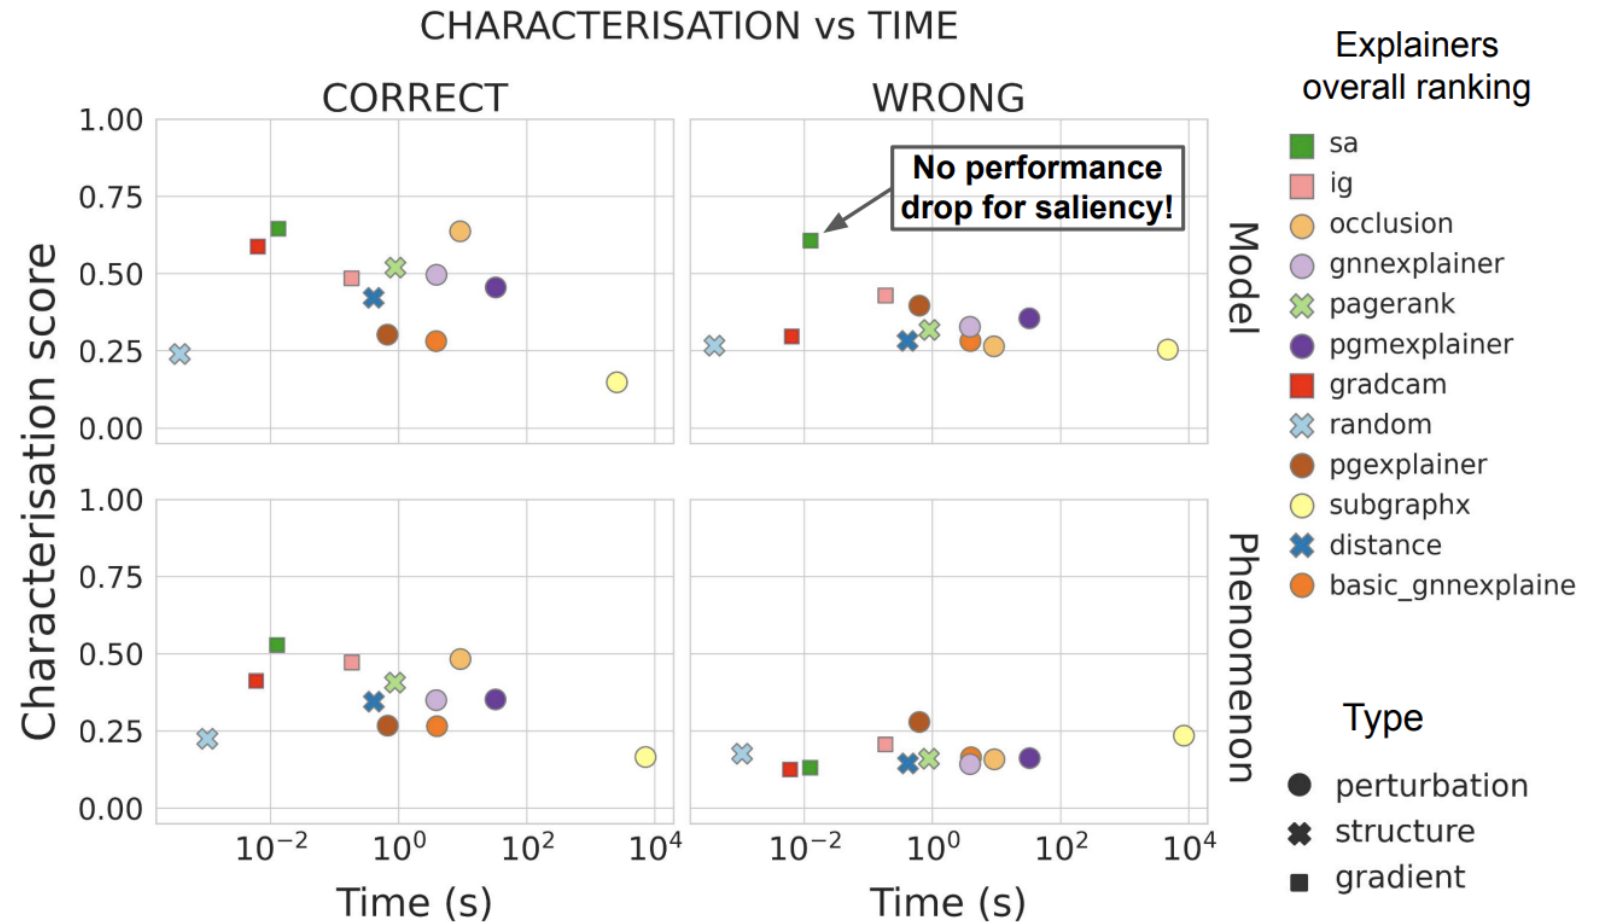
# Results: Explain Efficiency vs. Characterization Score

- Saliency has the highest overall characterization score and efficiency.
- Occlusion has the best overall score in the setting of hard mask.



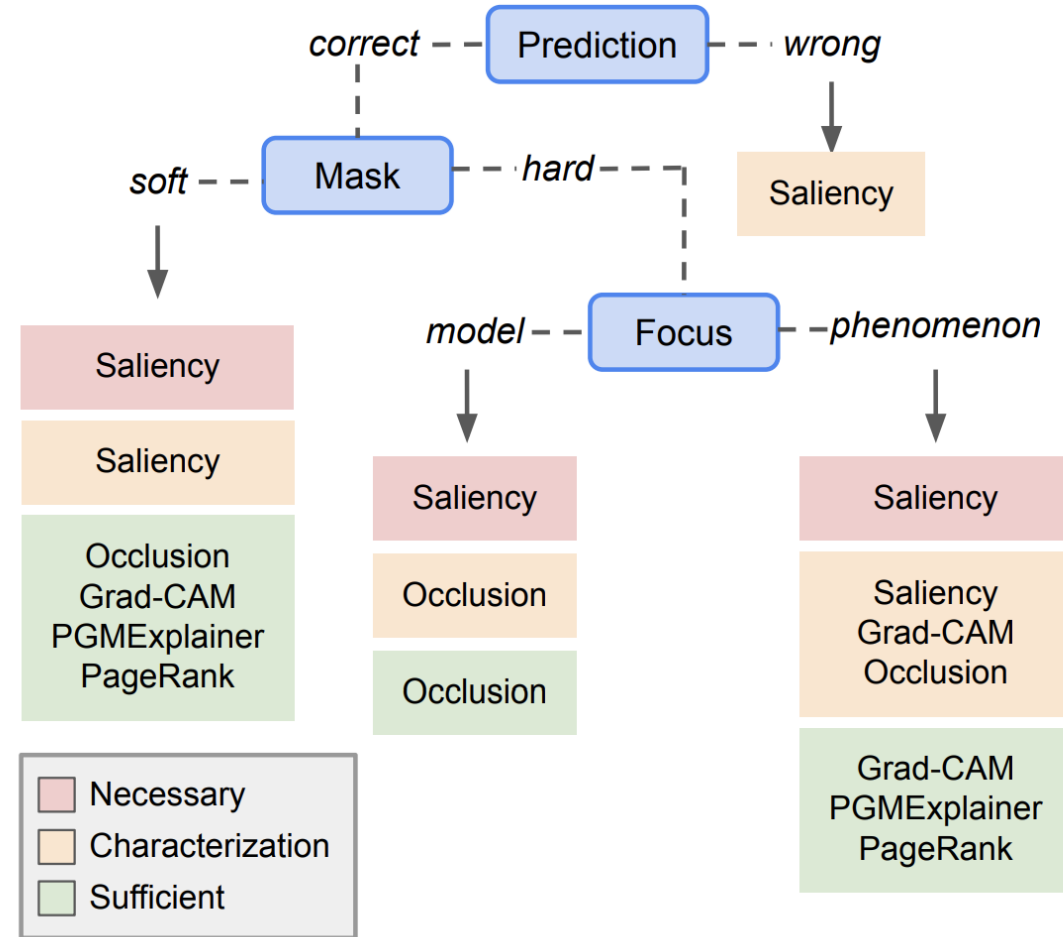
# Results: Correct & Wrong Predictions

- Correct predictions → phenomenon & model
- Wrong predictions → model



# Results: How to Select an Explainability Method

- Decision Tree which guides users to select the optimal method



# Summary of the Lecture

- Trustworthy GNN
  - Robustness, explainability, privacy, fairness, accountability, environmental well-being,...
- Adversarial Attacks and Robustness of GNNs
  - Adversarial examples
  - Attack possibilities: direct attack, indirect attack
  - Mathematical formulation
- Explainability of GNN
  - GNNExplainer
  - GraphFramEx